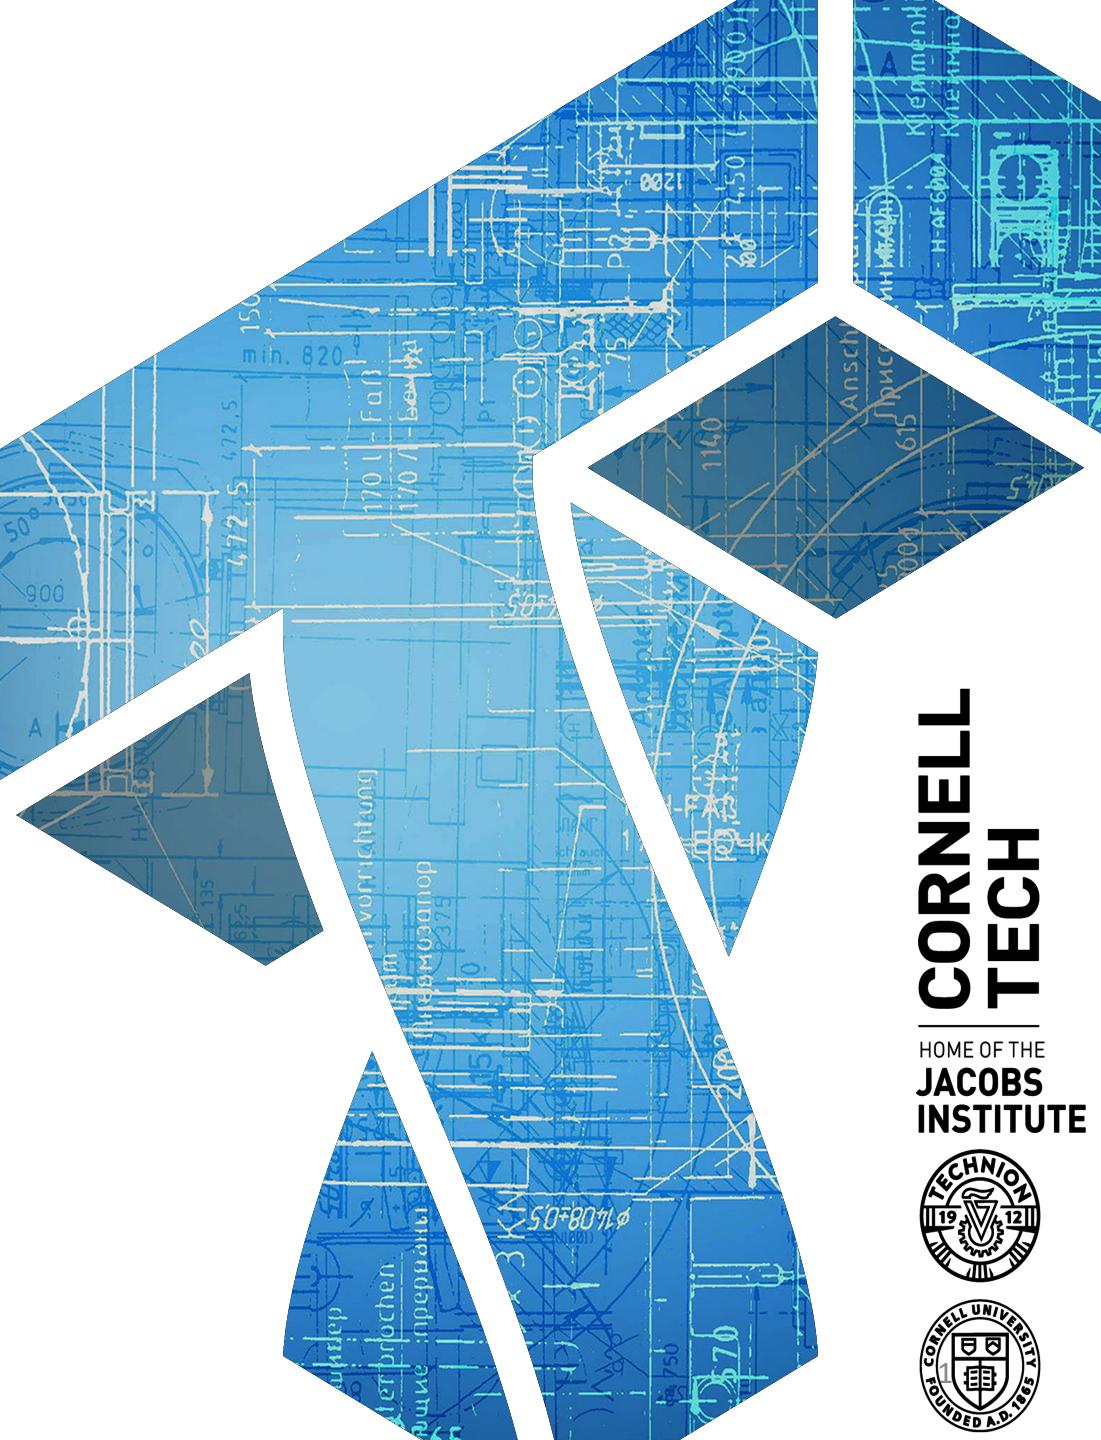


# CS 6431: Padding oracles

Instructor: Tom Ristenpart

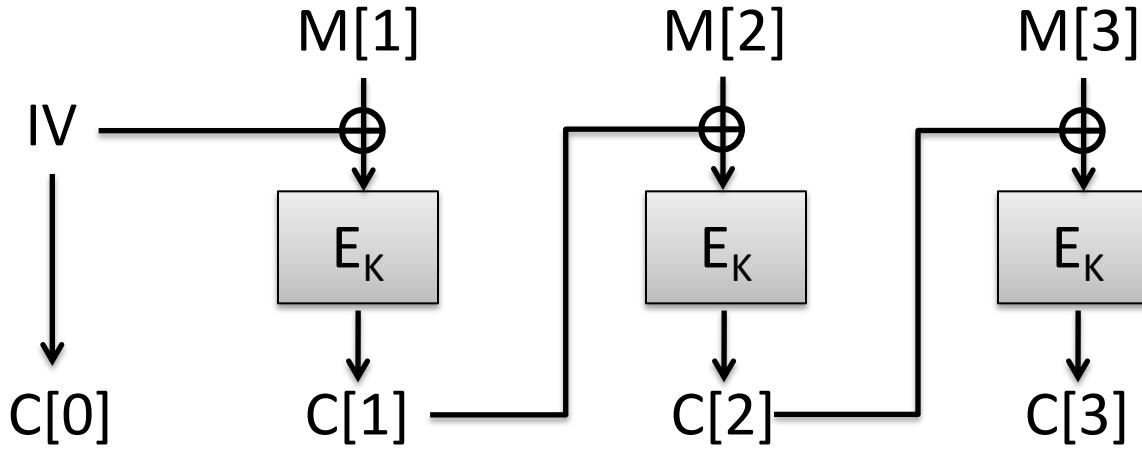
<https://github.com/tomrist/cs6431-fall2021>



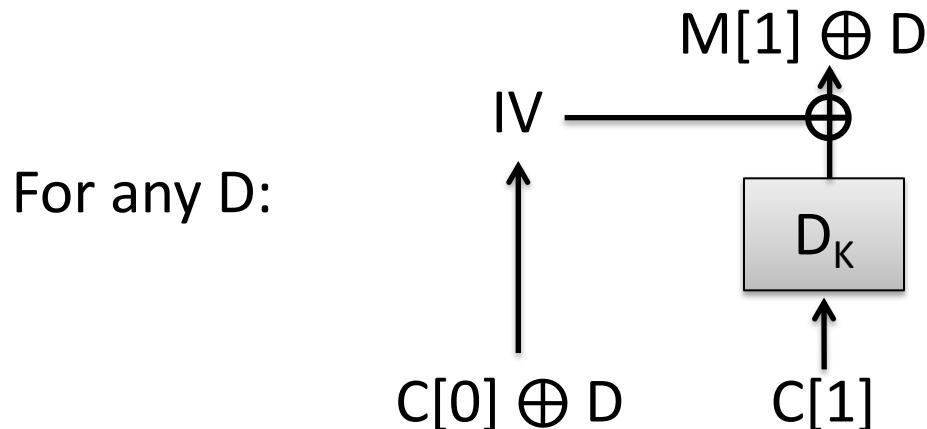
# Proposal & reviews being reviewed

- So far nice job!
  - A few missing reviews this morning. You will be docked for late reviews, reach out to explain please.
- Feel free to discuss proposals/reviews on HotCRP
  - Get clarification from reviewers
  - Provide clarification to reviewers
- I hope to finish going over all the project proposals by the end of the weekend. Will upload reviews
- Questions?

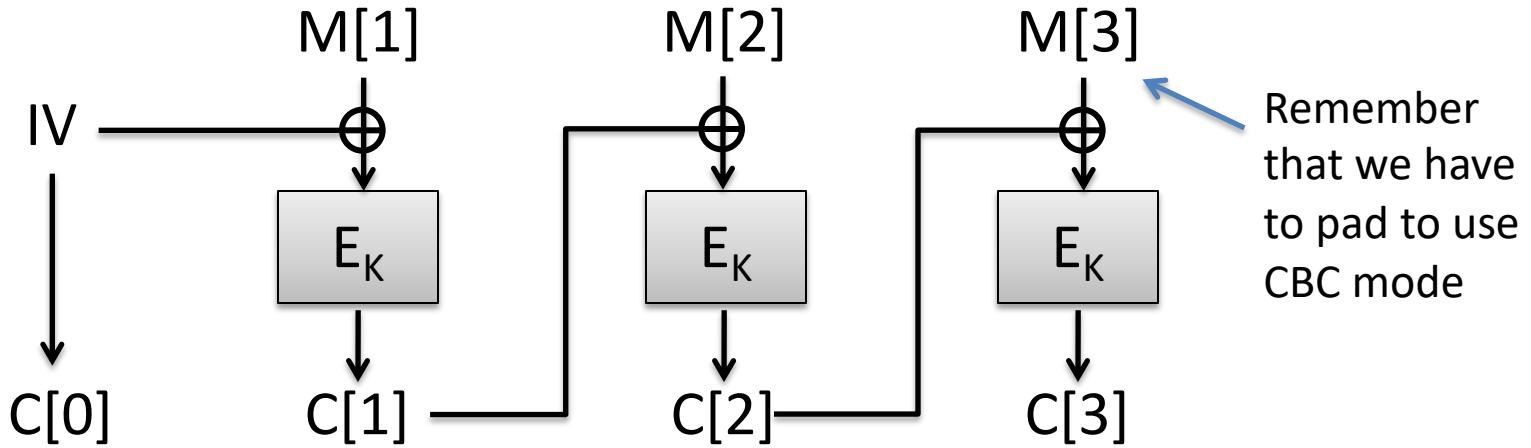
# CBC mode is malleable



How do we change bits of M received by server in controlled way?



# Padding oracle attacks



Padding oracle attacks take advantage of decryption implementations that reveal when padding checks fail

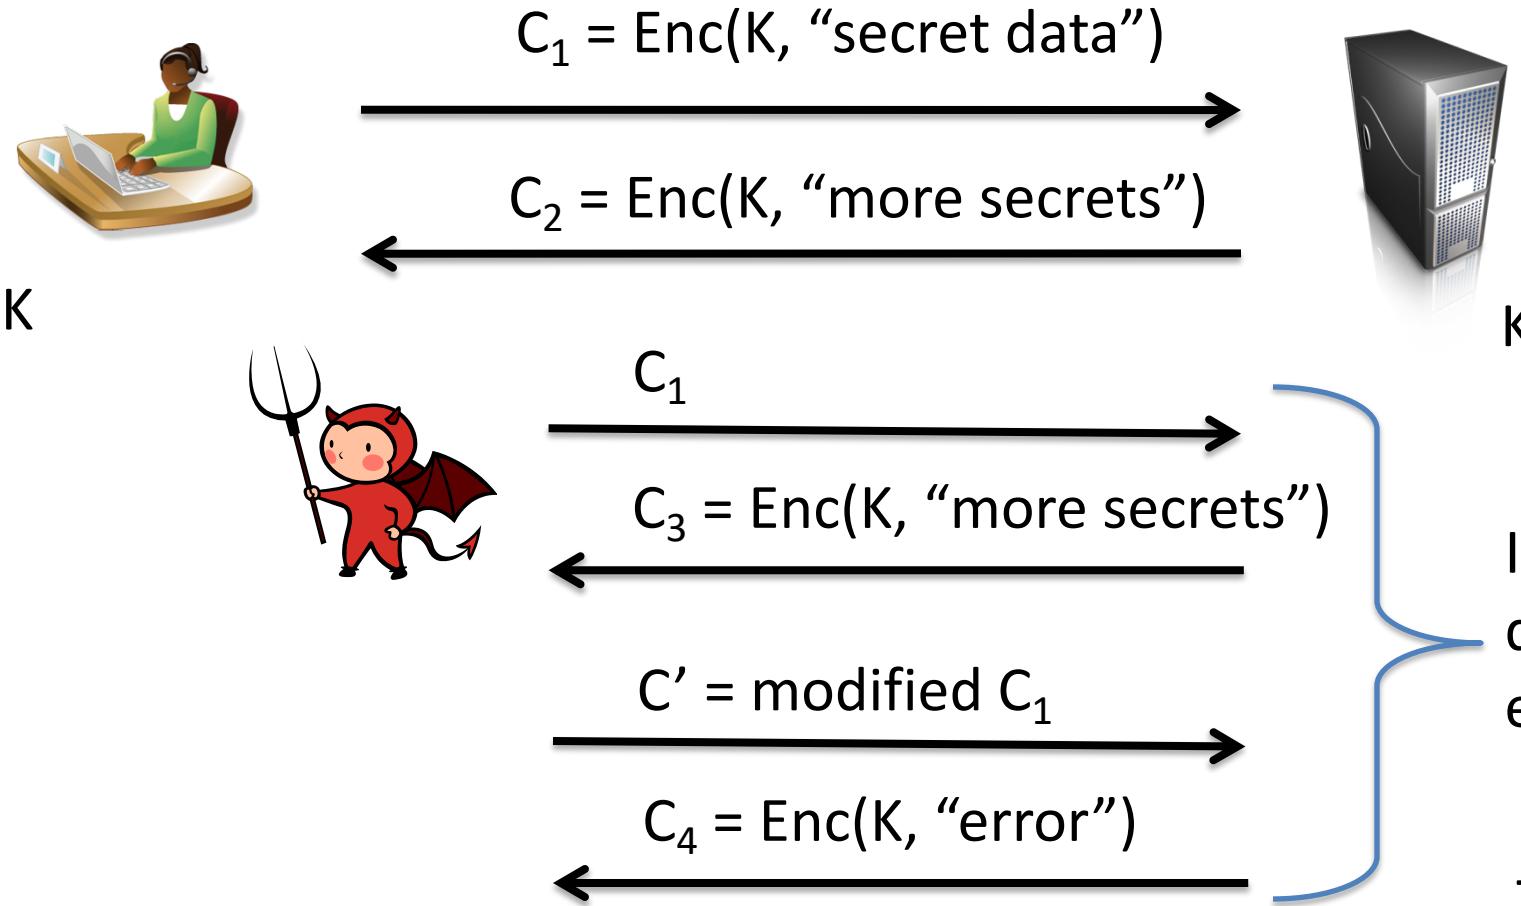
In most cases this allows *complete recovery of plaintexts*, violating confidentiality

Widespread damage: TLS, ASP.net, XML encryption, ...

# Partial decryption oracle threat model

- Attacker obtains encrypted message(s)  $\text{enc}(K, M)$  for some unknown  $M$
- Attacker wants to recover (information about  $M$ )
  - Can query a *partial decryption oracle* that does not reveal  $M$ , but some limited information about decryption process

# Partial decryption oracles arise frequently in practice



K

K

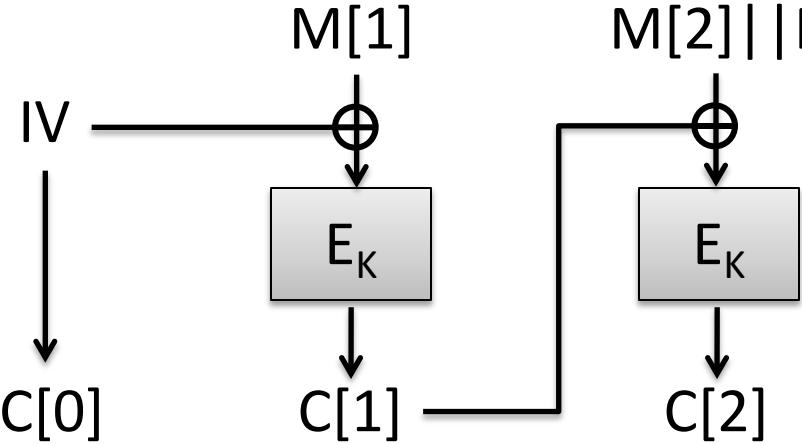
In practice usually easy to distinguish  $C_3$  from  $C_4$  even without K

$|C_4| \neq |C_3|$

Timing differs for successful vs. unsuccessful decryption

TLS/HTTPS canonical examples where decryption oracles arise

# Simple situation: pad by 1 byte



Assume that  
 $M[1] \parallel M[2]$  has length  
2n-8 bits

P is one byte of padding  
that must equal 0x00



Adversary  
obtains  
ciphertext  
 $C[0], C[1], C[2]$

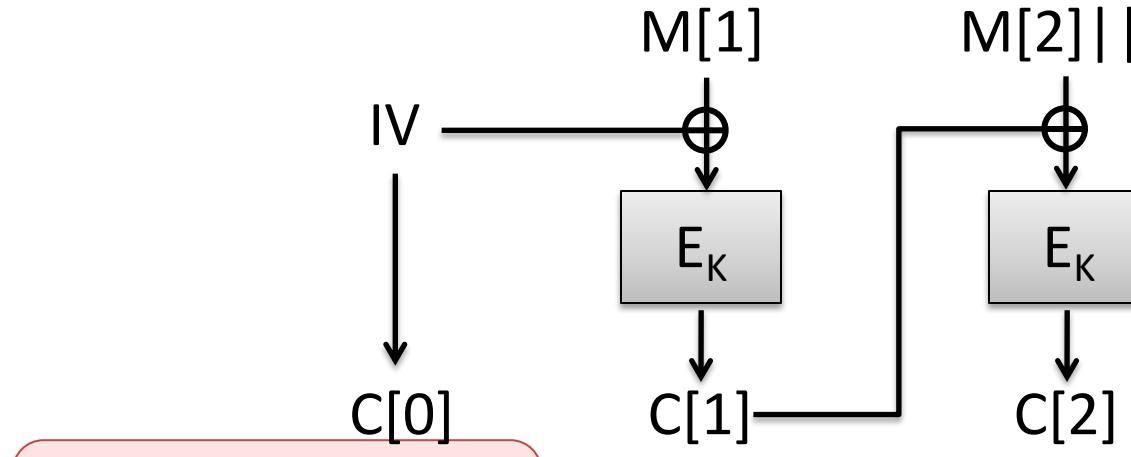
$C[0], C[1], C[2]$   
ok

$C[0], C[1] \oplus 1, C[2]$   
error



Dec( $K, C'$ )  
 $M'[1] \parallel M'[2] \parallel P' = \text{CBC-Dec}(K, C')$   
If  $P' \neq 0x00$  then  
Return error  
Else  
Return ok

# Simple situation: pad by 1 byte



Low byte of M1  
equals i



Adversary  
obtains  
ciphertext  
 $C[0], C[1], C[2]$   
Let R be arbitrary  
n bits

$R, C[0], C[1]$

error

$R, C[0] \oplus 1, C[1]$

error

$R, C[0] \oplus 2, C[1]$

error

...

$R, C[0] \oplus i, C[1]$

ok

Assume that  
 $M[1] || M[2]$  has length  
2n-8 bits

P is one byte of padding  
that must equal 0x00



```
Dec(K, C')
M'[1] || M'[2] || P' = CBC-Dec(K, C')
If P' ≠ 0x00 then
    Return error
Else
    Return ok
```

# PKCS #7 padding used most often with CBC mode

$$\text{PKCS#7-Pad}(M) = M || \underbrace{P || P || \dots || P}_{P \text{ repetitions of byte encoding number of bytes padded}}$$

Possible paddings:

01

02 02

03 03 03

04 04 04 04

...

FF FF FF FF ... FF

For block length of 16 bytes, don't need more than 16 bytes  
of padding (10 10 ... 10)

# Decryption (assuming at most one block of padding)

Dec( K, C )

$M'[1] \parallel \dots \parallel M'[m] = \text{CBC-Dec}(K, C)$

$P = \text{RemoveLastByte}(M'[m])$

For  $\text{ctr} < \text{int}(P)$ :

$P' = \text{RemoveLastByte}(M'[m])$

If  $P' \neq P$  then

Return error

$\text{ctr}++$

Return ok

After operation,  
 $M'[m]$  is one byte shorter

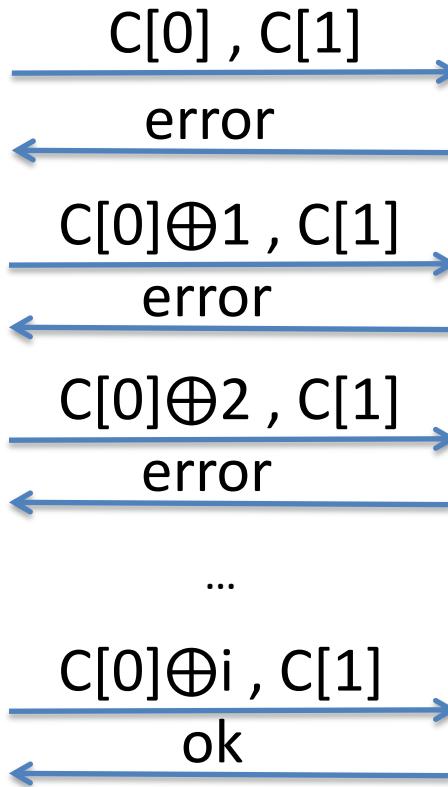
- “ok” / “error” stand-ins for some other behavior:
- Passing data to application layer (web server)
  - Returning other error code (if padding fails)

# PKCS #7 padding oracles

Low byte of  $M[1]$  most likely equals  $i \oplus 01$



Adversary obtains ciphertext  $C[0], C[1], C[2]$



Dec( K, C )

$M'[1] || \dots || M'[m] = \text{CBC-Dec}(K, C)$

$P = \text{RemoveLastByte}(M'[m])$

while  $\text{ctr} < \text{int}(P)$ :

$P' = \text{RemoveLastByte}(M'[m])$

If  $P' \neq P$  then

Return error

$\text{ctr}++$

Return ok

Why? Let  $X[1] = D(K, C_1)$

$$C[0][16] \oplus X[1][16] = M[1][16]$$

$$C[0][16] \oplus i \oplus X[1][16] = 01$$

$$M[1][16] \oplus i = 01$$

Actually, it could be that:

$$M[1][16] \oplus i = 02$$

Implies that  $M[1][15] = 02$

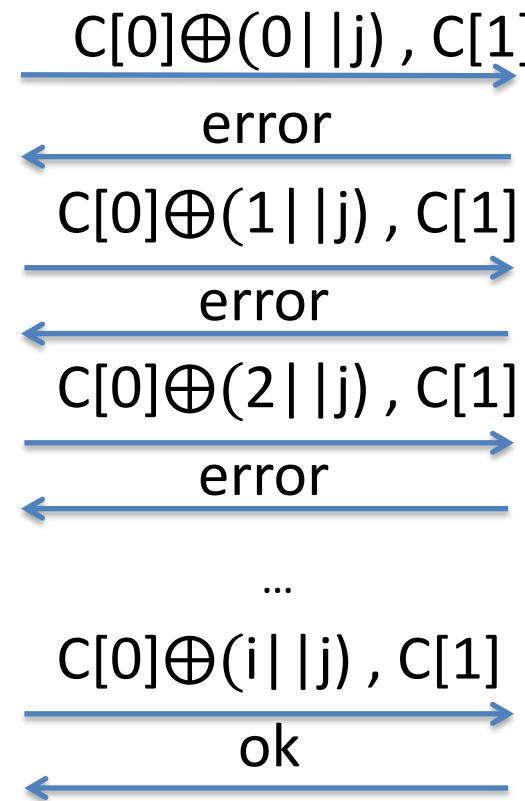
We can rule out with an additional query

# PKCS #7 padding oracles

Second lowest byte of  
M[1] equals  $i \oplus 02$



Adversary  
obtains  
ciphertext  
 $C[0], C[1], C[2]$



Dec( K, C )

$M'[1] \parallel \dots \parallel M'[m] = \text{CBC-Dec}(K, C)$

$P = \text{RemoveLastByte}(M'[m])$

while  $\text{ctr} < \text{int}(P)$ :

$P' = \text{RemoveLastByte}(M'[m])$

If  $P' \neq P$  then

Return error

$\text{ctr}++$

Return ok

Set  $j = M[1][16] \oplus 01 \oplus 02$

Keep going to recover entire block of message  $M[1]$   
Can repeat with other blocks  $M[2], M[3], \dots$   
Worst case:  $256 * 16$  queries per block

# Chosen ciphertext attacks against CBC

Attack	Description	Year
Vaudenay	10's of chosen ciphertexts, recovers message bits from a ciphertext. Called "padding oracle attack"	2001
Canvel et al.	Shows how to use Vaudenay's ideas against TLS	2003
Degabriele, Paterson	Breaks IPsec encryption-only mode	2006
Albrecht et al.	Plaintext recovery against SSH	2009
Duong, Rizzo	Breaking ASP.net encryption	2011
Jager, Somorovsky	XML encryption standard	2011
Duong, Rizzo	"Beast" attacks against TLS	2011
AlFardan, Paterson	Attack against DTLS	2012
AlFardan, Paterson	Lucky 13 attack against DTLS and TLS	2013
Albrecht, Paterson	Lucky microseconds against Amazon's s2n library	2016

# Duong, Rizzo padding oracle work

- WOOT 2010 paper:
  - Using padding oracles to build ciphertext of arbitrary plaintext
  - Ways to find padding oracles in practice
  - Case study of problems
- Oakland 2011 paper:
  - Focus on ASP.net, a critical web framework by Microsoft

# Session handling and login



GET /index.html

Set-Cookie: AnonSessID=134fds1431



Protocol  
is HTTPS.  
Elsewhere  
else just HTTP.

Nowadays  
increasingly all  
HTTPS

POST /login.html?name=bob&pw=12345

Cookie: AnonSessID=134fds1431

Set-Cookie: SessID=83431Adf

GET /account.html

Cookie: SessID=83431Adf

# Security problems here?



website.com



POST /login.html?name=bob&pw=12345

Cookie: AnonSessID=134fds1431



Set-Cookie: SessID=83431Adf

Secret key K only  
known to server

GET /account.html

Cookie: SessID=83431Adf

83431Adf = Enc(K, "admin=0")

Malicious client can simply flip a few bits to change admin=1

Example of an ***integrity / authenticity violation***

Duong-Rizzo show how to turn into padding oracle, do all sorts of damage

# Duong, Rizzo padding oracle work

- WOOT 2010 paper:
  - Using padding oracles to build ciphertext of arbitrary plaintext
  - Ways to find padding oracles in practice
  - Case study of problems
- Oakland 2011 paper:
  - Focus on ASP.net, a critical web framework by Microsoft
  - “The best attack requires less than 2,000 HTTP requests to steal secret cryptographic keys in ASP.NET applications. After obtaining those keys, the attacker can easily create authentication tickets that he can use to sign in to any user account.”

# Security problems here?



POST /WebResource.axd?d=encrypted\_id&t=timestamp

HTTP 500 or HTTP 404

ASP.net website



Secret key K only known to server

encrypted\_id = CBC-Enc(K, R#anything || ~/path/to/file)

---

**Algorithm 3** Downloading files using CBC-R and the padding oracle in WebResource.axd.

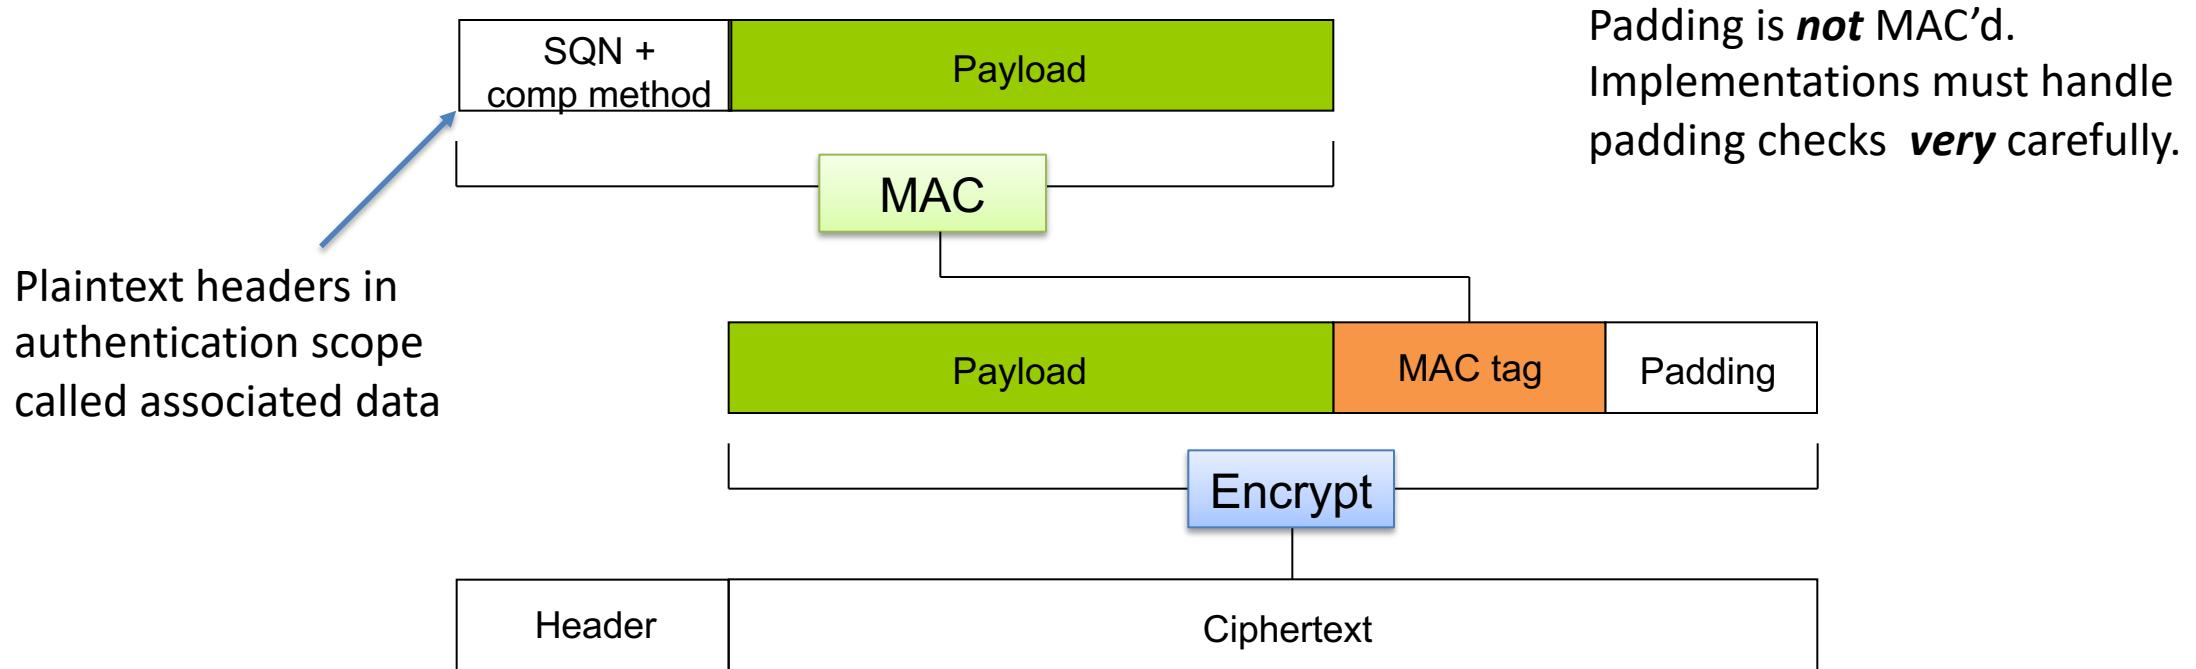
---

- 1) Pad the message  $\|\|~\text{/path/to/file}$  according to the PKCS#5 padding scheme, and set the whole result as the target  $P$ .
  - 2) Divide  $P$  into  $n$  blocks of  $b$  bytes, denote them  $P_1, P_2, \dots, P_n$ .
  - 3) Use CBC-R to build  $C_n, \dots, C_2, C_1$  so that  $C_1|C_2|\dots|C_n$  decrypts to  $P_{garbage}|P_1|P_2|\dots|P_n$ .
  - 4) Pick a few random bytes  $r_1, r_2, \dots, r_b$ , and set  $C_0 = r_1|r_2|\dots|r_b$ .
  - 5) Set  $d = C_0|C_1|\dots|C_n$ , and send it to ScriptResource.axd to see if  $\text{/path/to/file}$  is downloaded. If no, go back to step 4.
  - 6) Output  $d$ .
-

# Authenticated encryption

- Malleability problems and chosen ciphertext attacks motivate *authenticated* encryption
  - No computationally bound adversary can forge a ciphertext that decrypts w/o error
- Actually authenticated encryption was already in place when Vaudenay's paper released

# TLS 1.2 record protocol: MAC-Encode-Encrypt (MEE)



The MAC (message authentication code) can be thought of as a pseudorandom function (PRF). During decryption:

Re-run MAC over header + Payload, check the given value

The hope is that this prevents manipulation of ciphertexts by adversaries, preventing padding oracle attacks

# (Simplified) Early TLS decryption

TLS-Decrypt(K, C)

$M'[1] \parallel \dots \parallel M'[m] = \text{CBC-Dec}(K, C)$

$P = \text{RemoveLastByte}(M'[m])$

while  $\text{ctr} < \text{int}(P)$ :

$P' = \text{RemoveLastByte}(M'[m])$

    If  $\text{plen}' \neq \text{plen}$

        Send PADDING\_ERROR message

$\text{ctr}++$

$(X, \text{tag}) \leftarrow \text{ParseLast20bytes}(M'[1] \parallel \dots \parallel M'[m])$

    If  $\text{MAC}(K, X) \neq \text{tag}$  then

        Send MAC\_ERROR MESSAGE

Return X as plaintext to application

How do we mount an attack?

# Authenticated encryption (w/ associated data)

- Encrypt-then-MAC generic composition
  - CTR mode encryption of message
  - HMAC or other MAC of ciphertext (including IV, associated data)
- Dedicated AEAD modes:
  - AES-GCM (Galois Counter Mode)
  - ChaCha2020/Poly1305
  - Offset Codebook mode (OCB)

# Evolution of views on symmetric encryption

1. Encryption is for confidentiality against passive adversaries
2. Add explicit message authentication layer to provide plaintext integrity
3. Need “all-in-one” authenticated encryption modes that simultaneously provide confidentiality and ciphertext integrity
  - Doesn’t say much if IV is chosen incorrectly, or if adversary has control over keys
4. Message franking [Facebook, Grubbs et al. 2017], partitioning oracle attacks [Len et al. 2020] show partial decryption oracle attacks against AES-GCM, other widely used schemes
5. Need authenticated encryption that is as robust as possible to other threat models:
  - IV (nonce) reuse, side channel resistance, key committing, ...