

CS 6431: Microarchitectural vulnerabilities

Instructor: Tom Ristenpart

<https://github.com/tomrist/cs6431-fall2021>



**CORNELL
TECH**

HOME OF THE
**JACOBS
INSTITUTE**



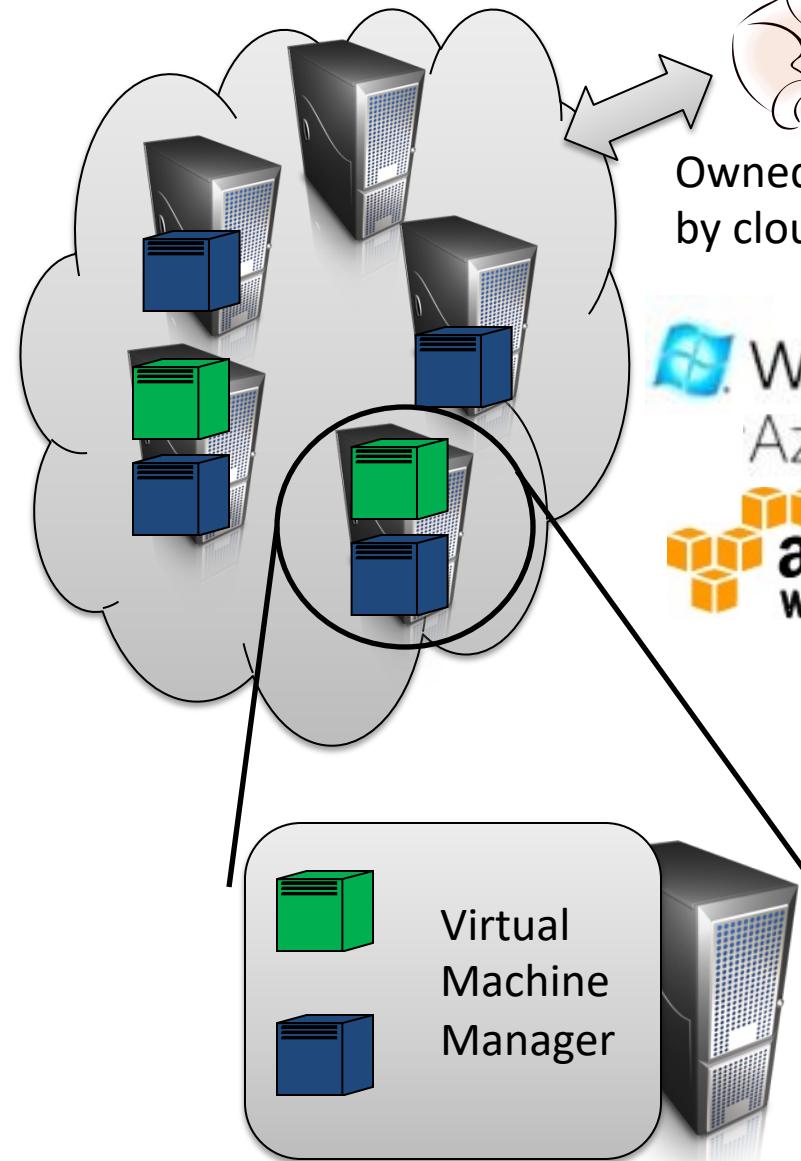
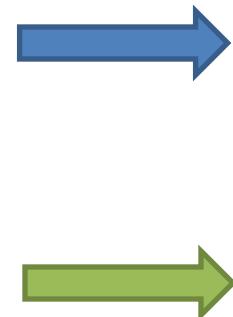
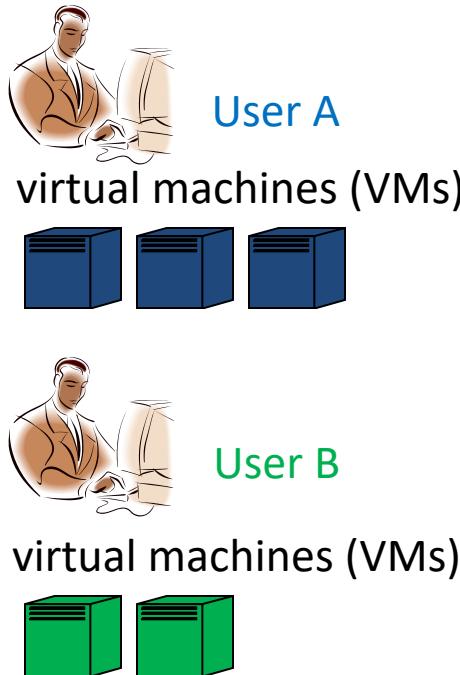
One motivating context: cloud computing

“If computers of the kind I have advocated become the computers of the future, then computing may someday be organized as a public utility just as the telephone system is a public utility... The computer utility could become the basis of a new and important industry.”

— John McCarthy, speaking at the MIT Centennial in 1961

A simplified model of public IaaS cloud computing

Users run Virtual Machines (VMs) on cloud provider's infrastructure



Owned/operated
by cloud provider

Windows Azure

amazon web services™



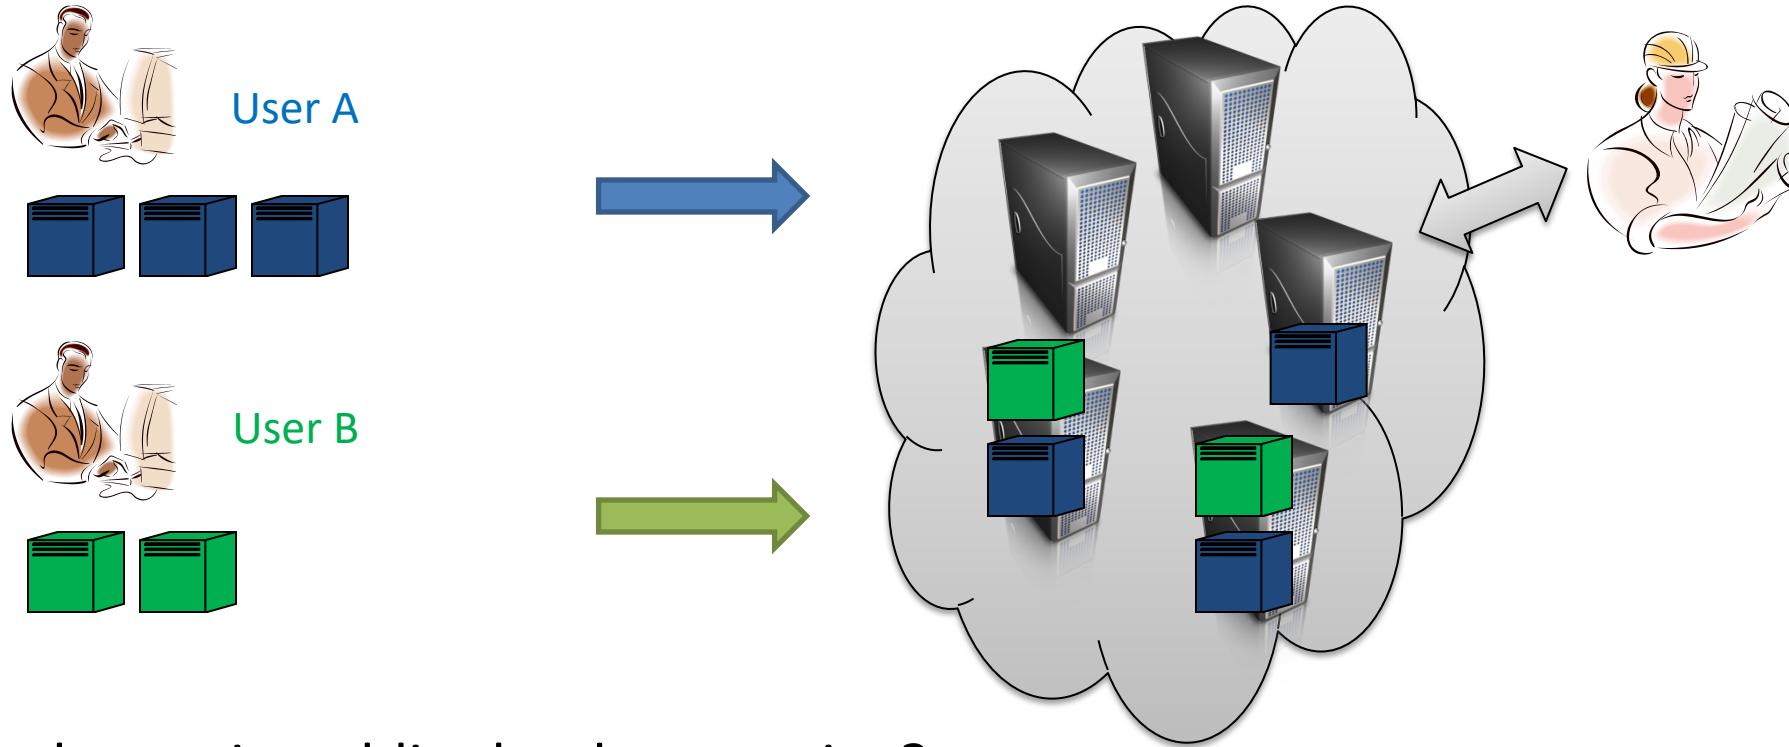
Google
Compute
Engine

Multitenancy (users share physical resources)

Virtual Machine Manager (VMM)
manages physical server resources for VMs

To the VM should look like dedicated server

Trust models in public cloud computing



Security threats in public cloud computing?

Provider spying on running VMs / data

External attacks against infrastructure

Cross-user attacks

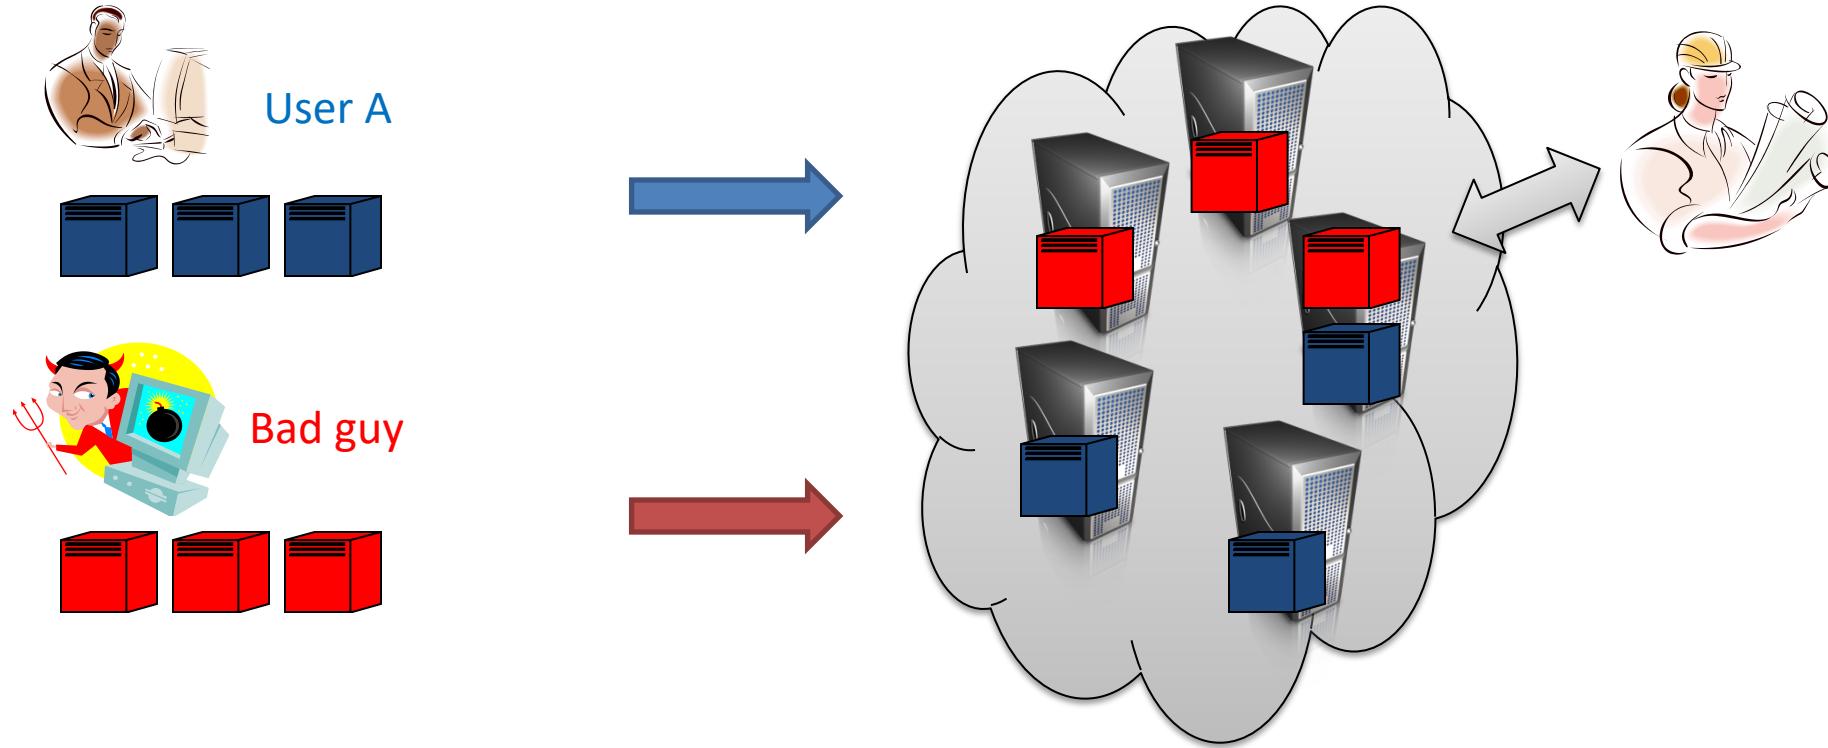
Escape-from-VM

Side-channels attacks

Degradation of service attacks

Resource-stealing attacks

Trust models in public cloud computing



Attacker identifies one or more victims VMs in cloud

1) Achieve advantageous placement via launching of VM instances

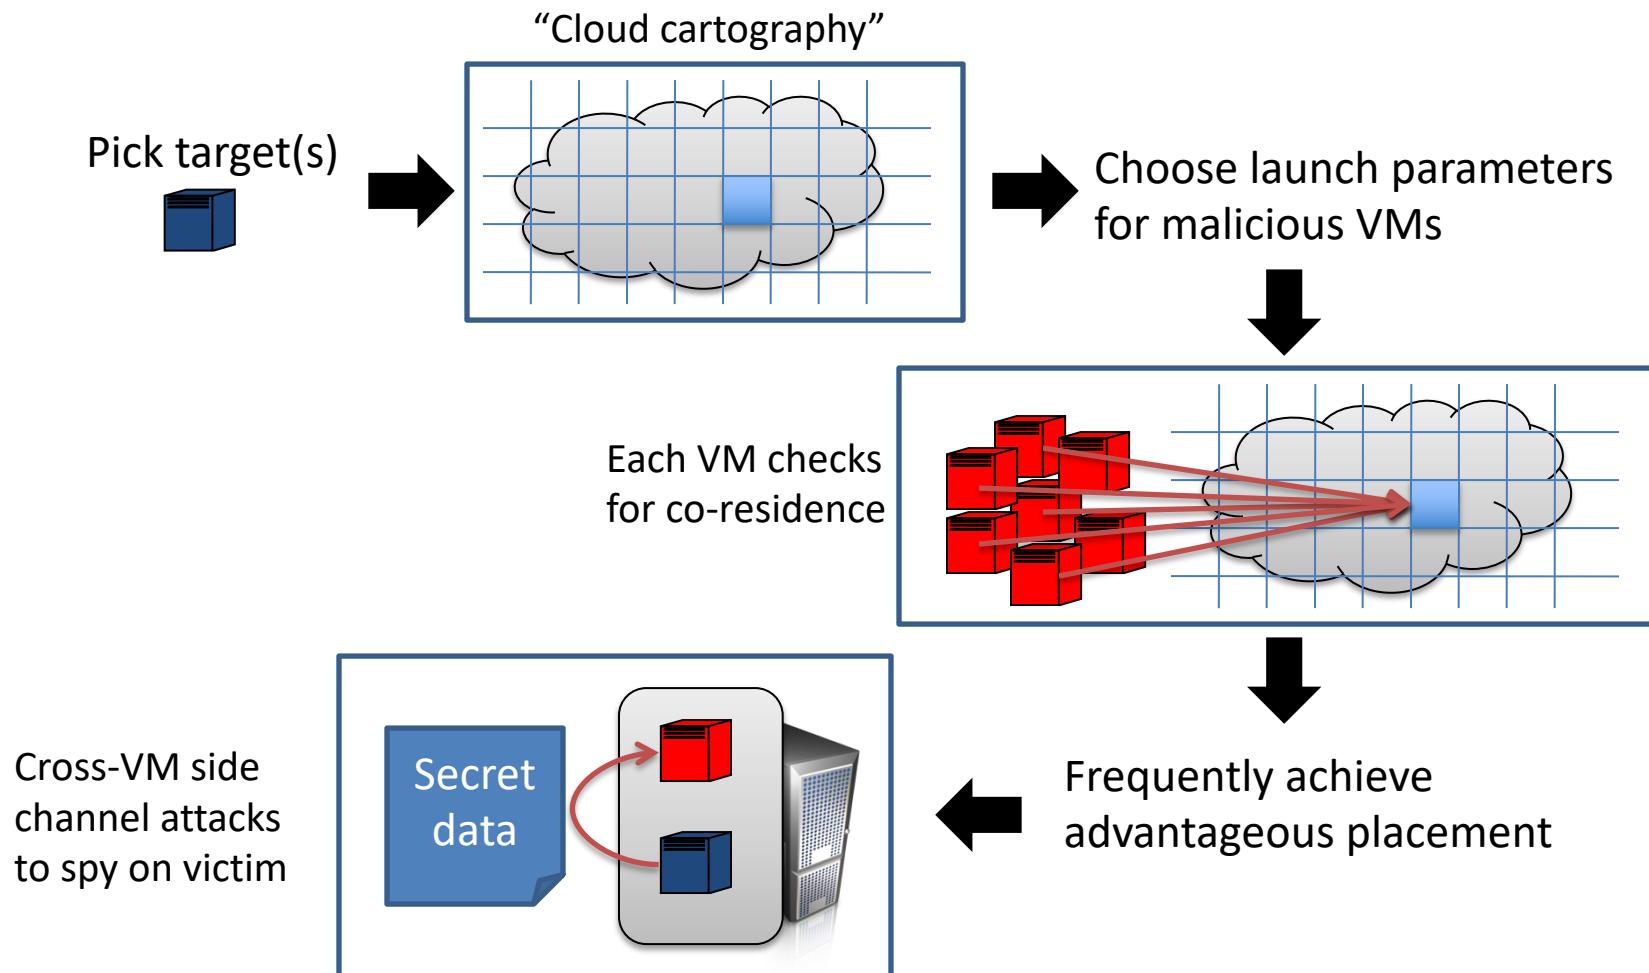
2) Launch attacks using physical proximity

Exploit VMM vulnerability

DoS

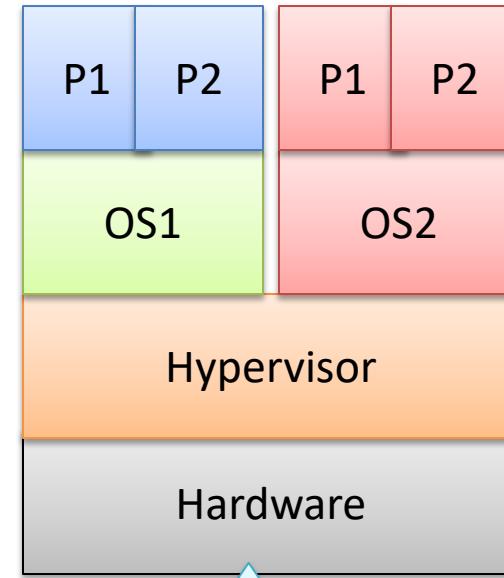
Side-channel attack

Adversary able to:



Violating isolation

- Covert channels between VMs circumvent access controls
 - Bugs in VMM
 - Side-effects of resource usage
- Degradation-of-Service attacks
 - Guests might maliciously contend for resources
 - Resource-stealing attacks (e.g., Xen scheduler vulnerability)
- Side channels
 - Spy on other guest via shared resources



Microarchitectural vulnerabilities:
- Side-channels
- Speculative execution bugs

Microarchitectural side-channels

- Lots of research on exploiting microarchitectural side-channels
 - CPU-cache based: Percivel 2005, Tromer et al. 2005
- Prime+Probe attacks [Tromer et al. 2005]
- Flush+Reload (F+R) more robust side-channel in shared memory settings [Yarom, Falkner 2013]



Client



Server

(Part of) TLS handshake for RSA transport

Pick random Nc

Check CERT
using CA public
verification key

Pick random PMS
 $C \leftarrow E(pk, PMS)$

ClientHello, MaxVer, Nc, Ciphers/CompMethods

ServerHello, Ver, Ns, SessionID, Cipher/CompMethod

CERT = (pk of bank, signature over it)

C

Pick random Ns

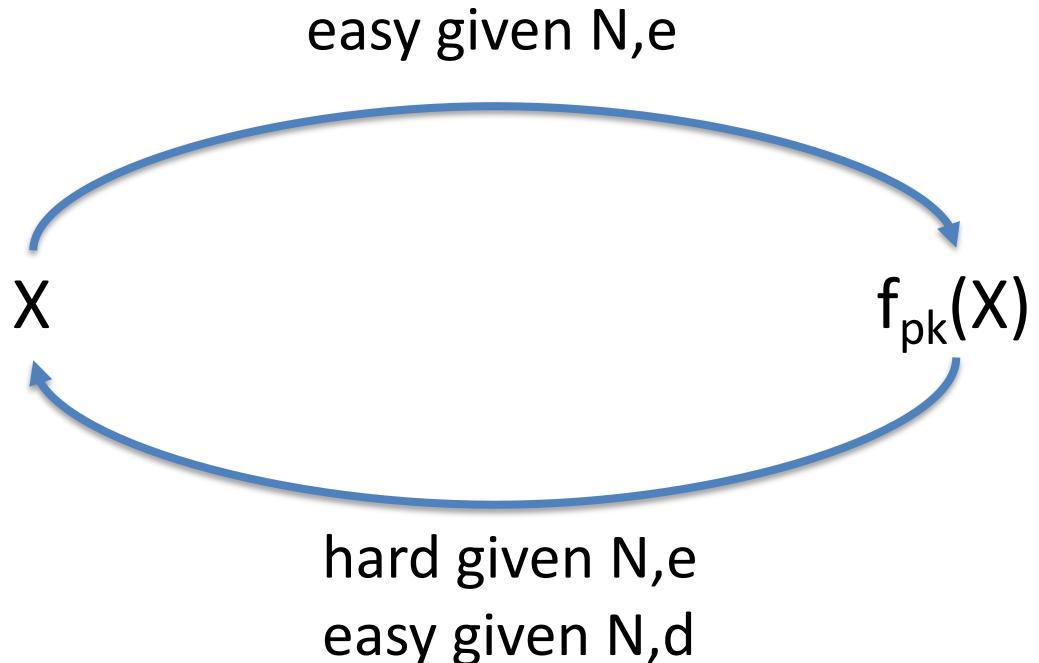
$PMS \leftarrow D(sk, C)$



The RSA trapdoor permutation

$$pk = (N, e) \quad sk = (N, d) \quad \text{with } ed \bmod \phi(N) = 1$$

$$f_{N,e}(x) = x^e \bmod N \quad g_{N,d}(y) = y^d \bmod N$$



```
SqrAndMulExp(y,d,N)
bk,...,b0 = d
f = 1
For i = k down to 0 do
    f = Square(y,N)
    If bi = 1 then
        f = Multiply(f,y,N)
Return f
```

SqrAndMulExp(y,d,N)

$b_k, \dots, b_0 = d$

$f = 1$

For $i = k$ down to 0 do

$f = \text{Square}(y, N)$

If $b_i = 1$ then

$f = \text{Multiply}(f, y, N)$

Return f

$$y^{11} = y^{8+2+1} = y^8 \bullet y^2 \bullet y$$

$$b_3 = 1 \quad f_3 = 1 \bullet y$$

$$b_2 = 0 \quad f_2 = y^2$$

$$b_1 = 1 \quad f_1 = (y^2)^2 \bullet y$$

$$b_0 = 1 \quad f_0 = (y^4 \bullet y)^2 \bullet y = y^8 \bullet y^2 \bullet y$$

$$d = \sum_{b_i \neq 0} 2^i$$

$$y^d \equiv y^{\sum_{b_i \neq 0} 2^i} \equiv \prod_{b_i \neq 0} y^{2^i} \pmod{N}$$

Highly vulnerable to side-channel attacks!

- Timing
- CPU State
- Power

$d_i = 1 \rightarrow \text{"SM"}$

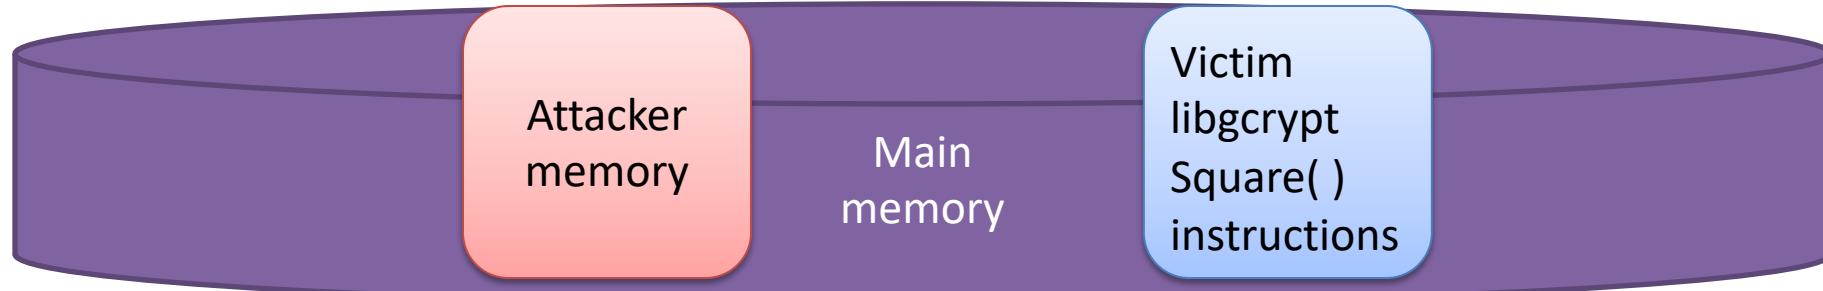
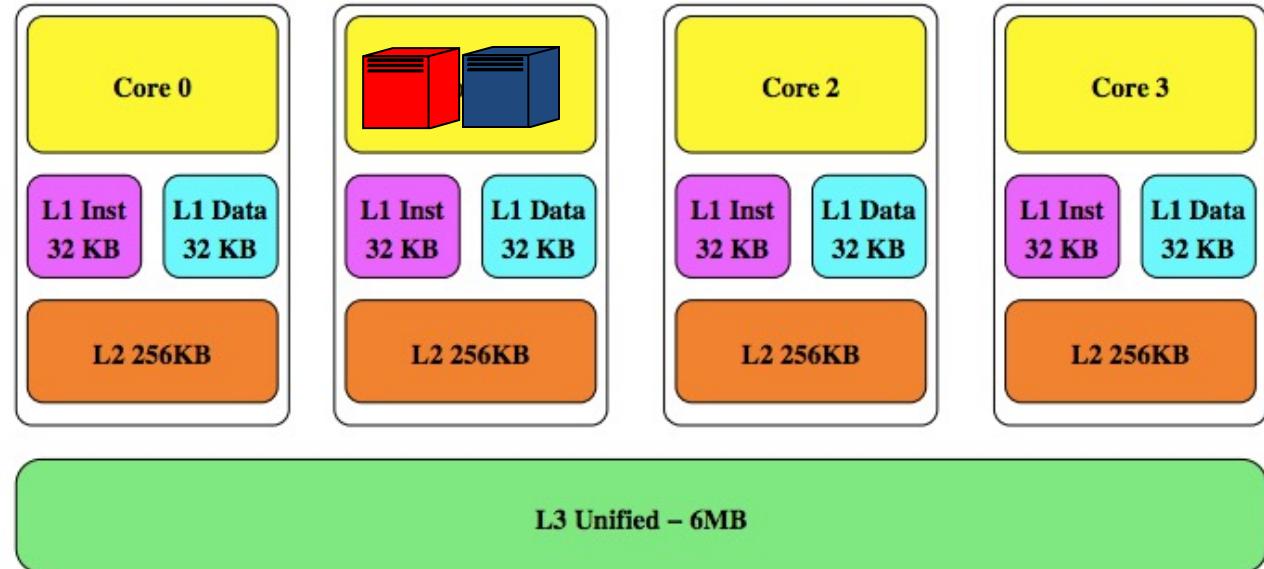
$d_i = 0 \rightarrow \text{"S"}$

Sequence of function calls
reveals secret key

Towards Prime+Probe

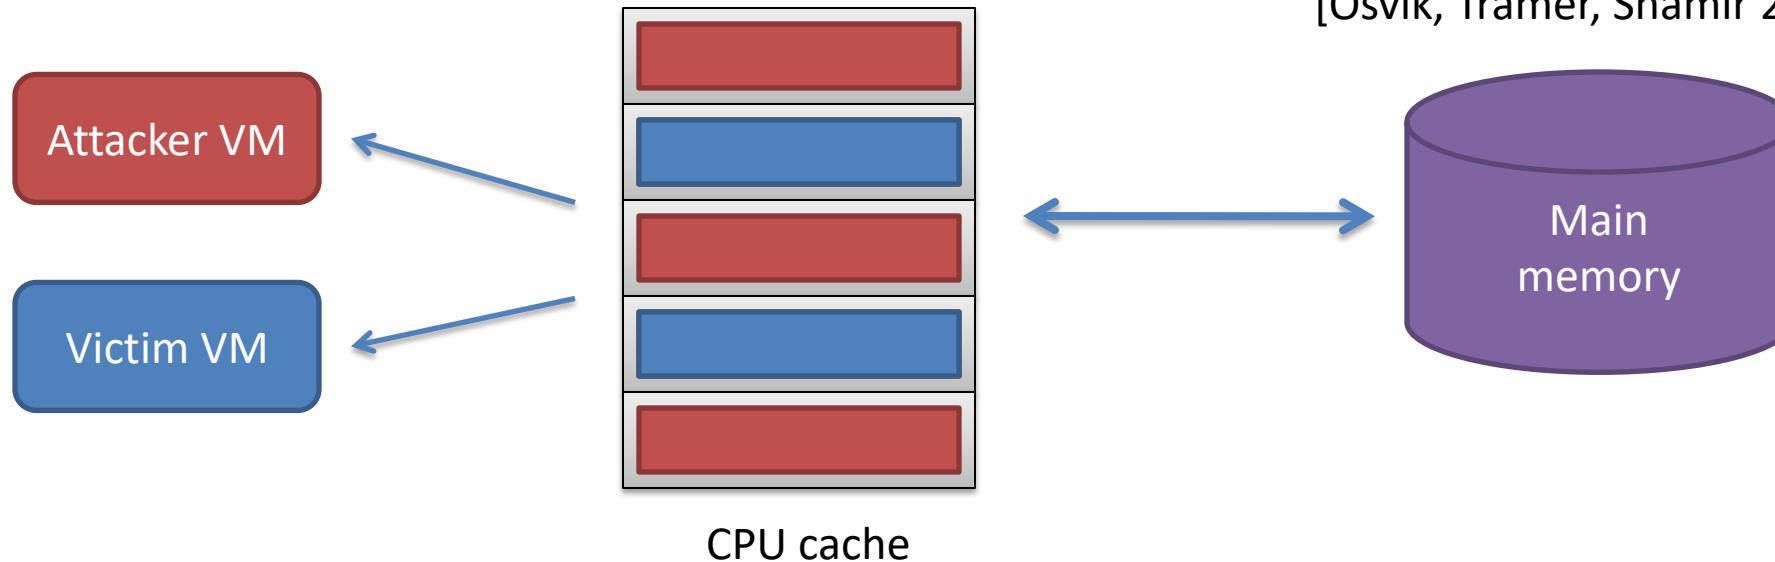
Suppose victim and attacker shares a core

Also sharing L1 instruction & data caches



Cross-VM side channels using CPU cache contention

Introduced in cross-process setting by [Osvik, Tramer, Shamir 2006]



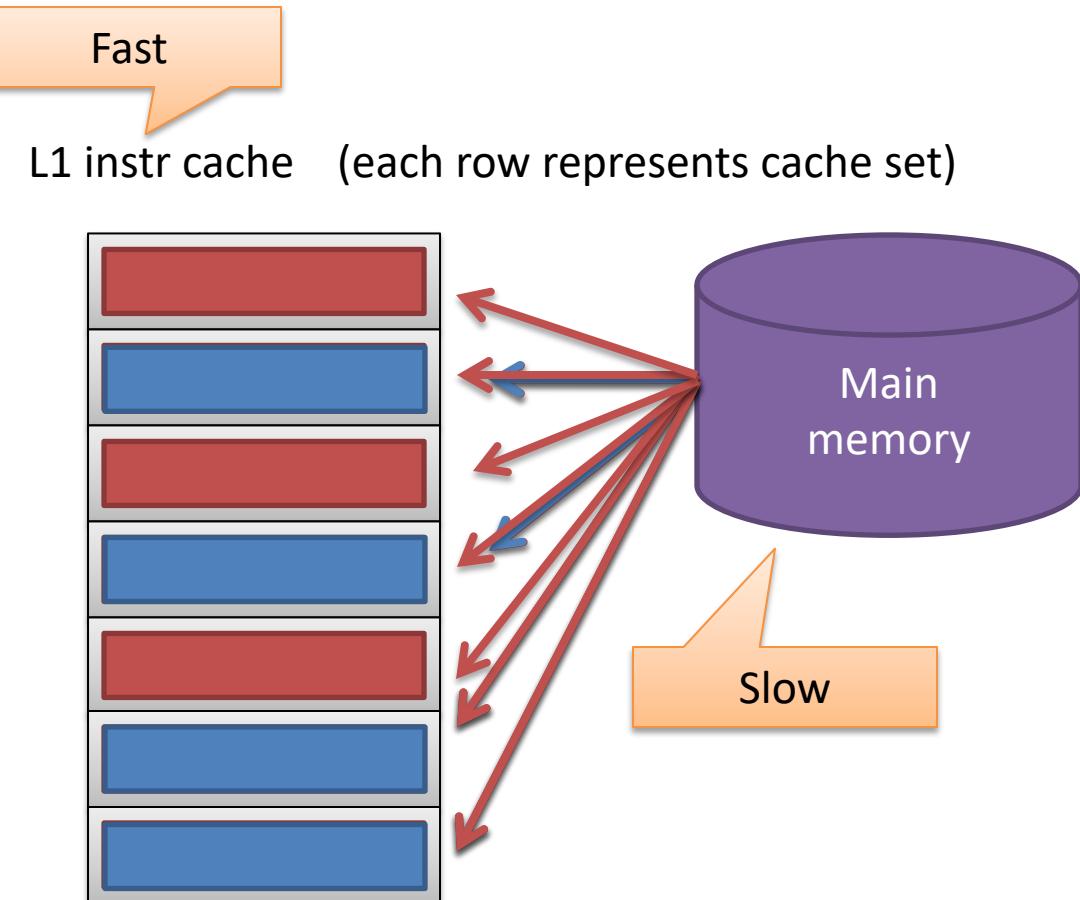
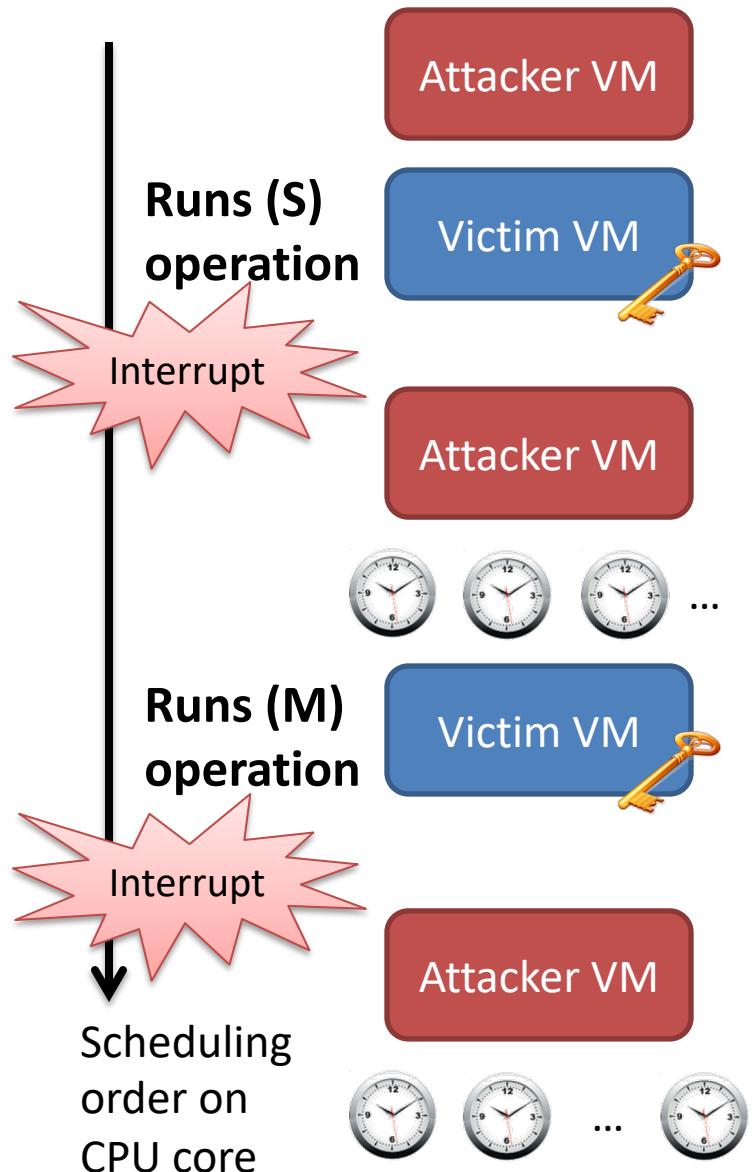
- 1) Read in a large array (fill CPU cache with attacker data)
- 2) Busy loop (allow victim to run)
- 3) Measure time to read large array (the load measurement)

Locations in cache occupied by victim will take longer to load



Information about victim's use of cache revealed to attacker

Prime+Probe protocol



- Timings correlated to (distinct) cache usage patterns of S, M operations
- Can spy frequently (every $\sim 16 \mu\text{s}$) by exploiting scheduling

Prime+Probe limitations

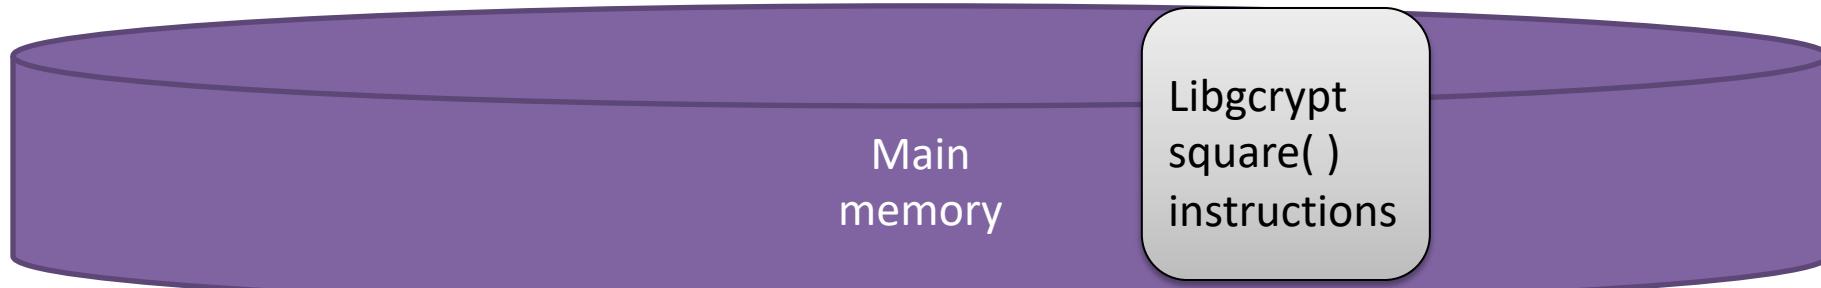
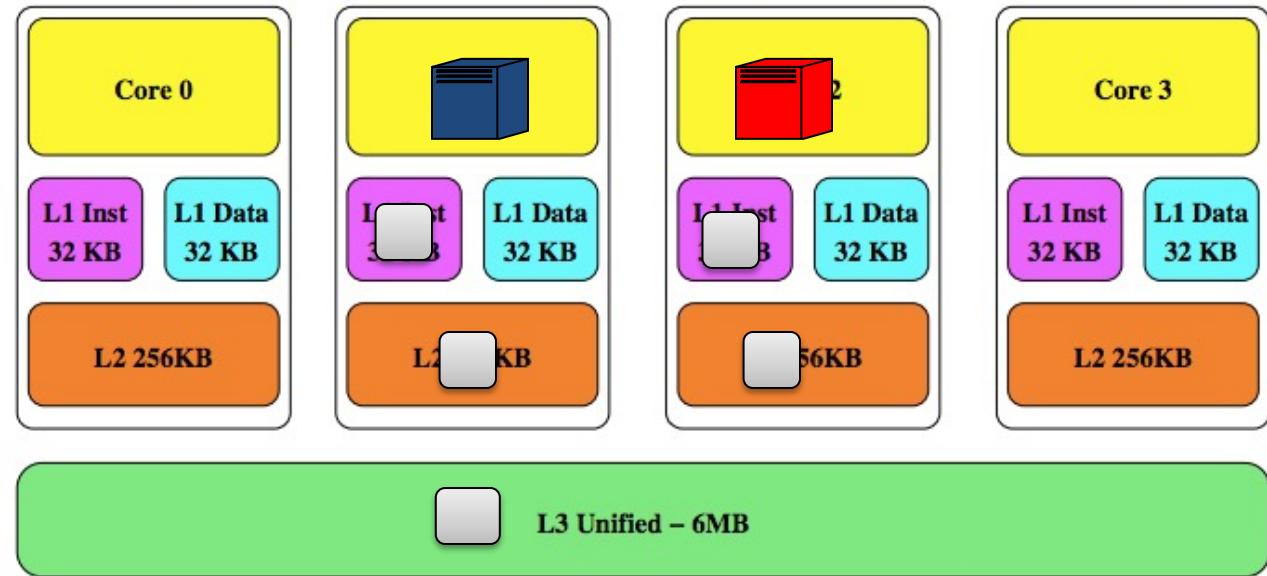
- Originally worked only for L1 caches
 - Some works extended to last-level cache (LLC) in certain settings
 - Multi-core settings difficult but feasible in lab (Zhang et al. 2012)
 - Inci et al. 2016 “Cache Attacks Enable Bulk Key Recovery on the Cloud”
- Lots of noise from various sources

Towards Flush+Reload

Deduplication-based
memory page sharing (Linux,
KVM, VMWare)

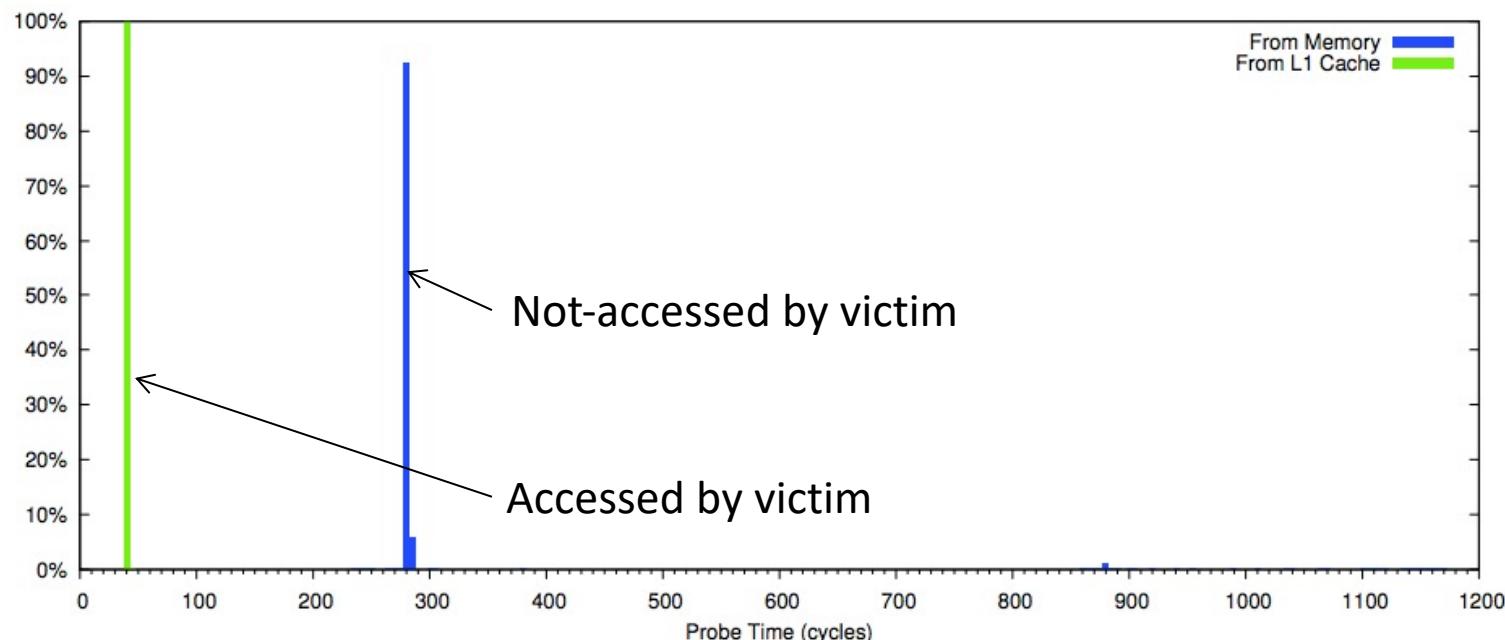
- Duplicate memory pages detected, physical pages coalesced
- Virtual address spaces different, but mapped to same physical addresses

Inclusive cache architecture

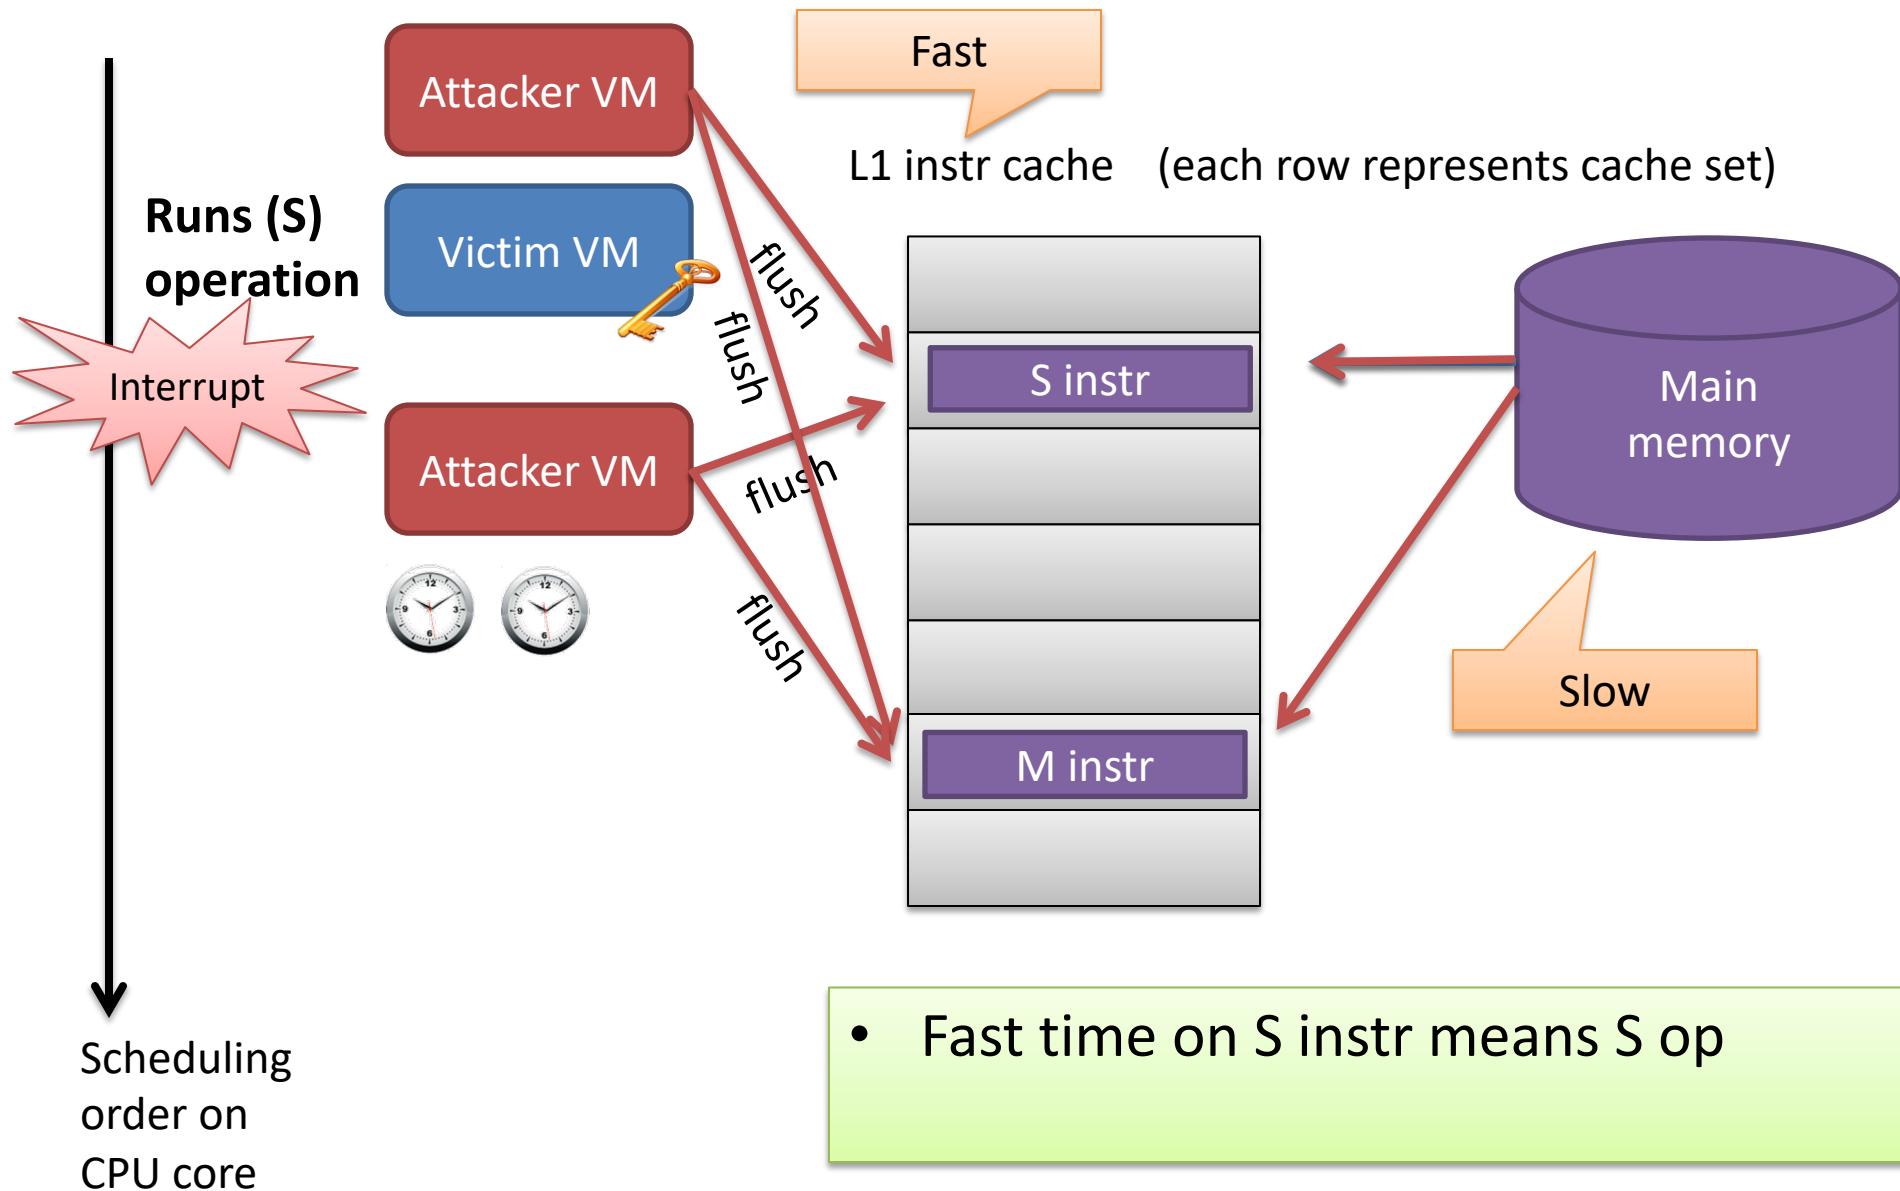


Flush+Reload protocol

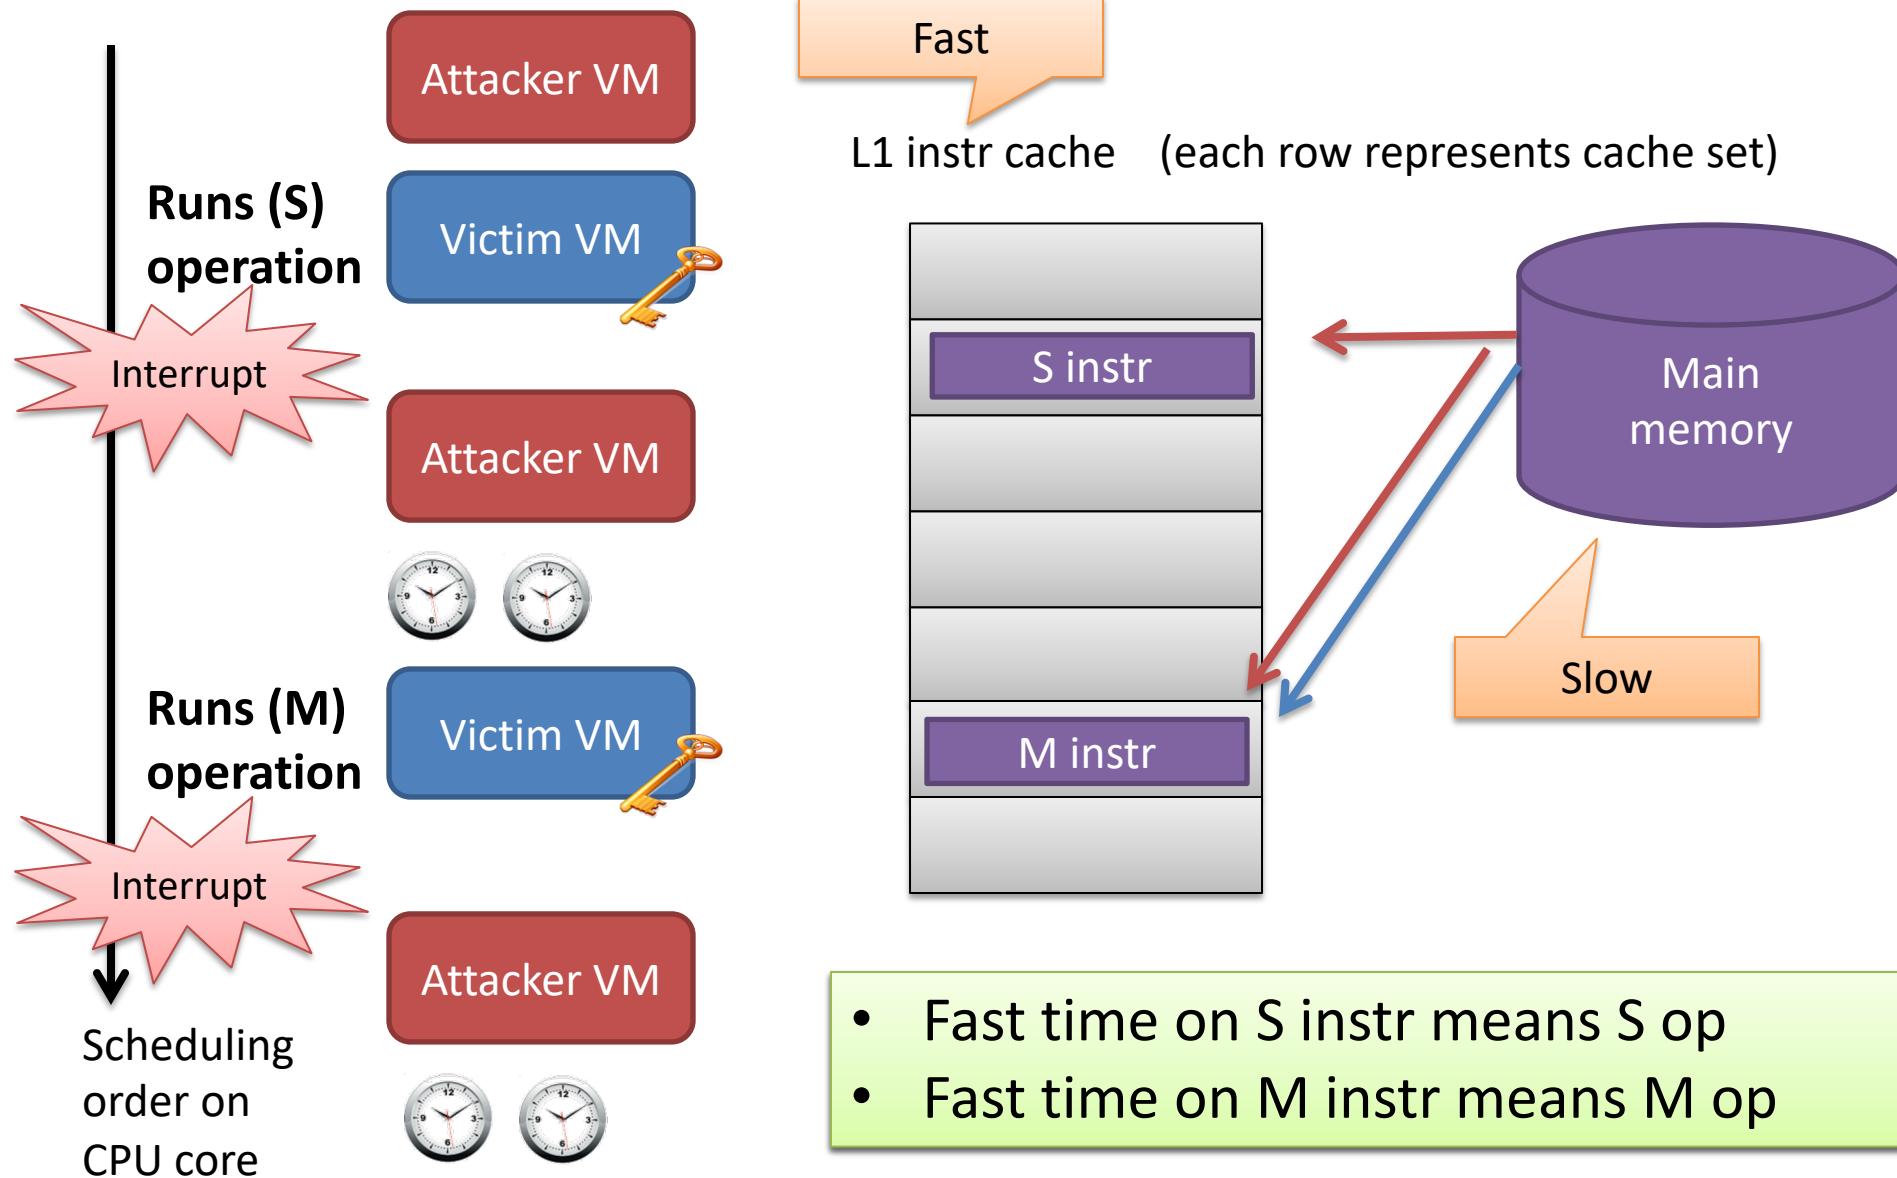
- Flush from LLC memory line of interest
- Wait
- Time reloading memory line



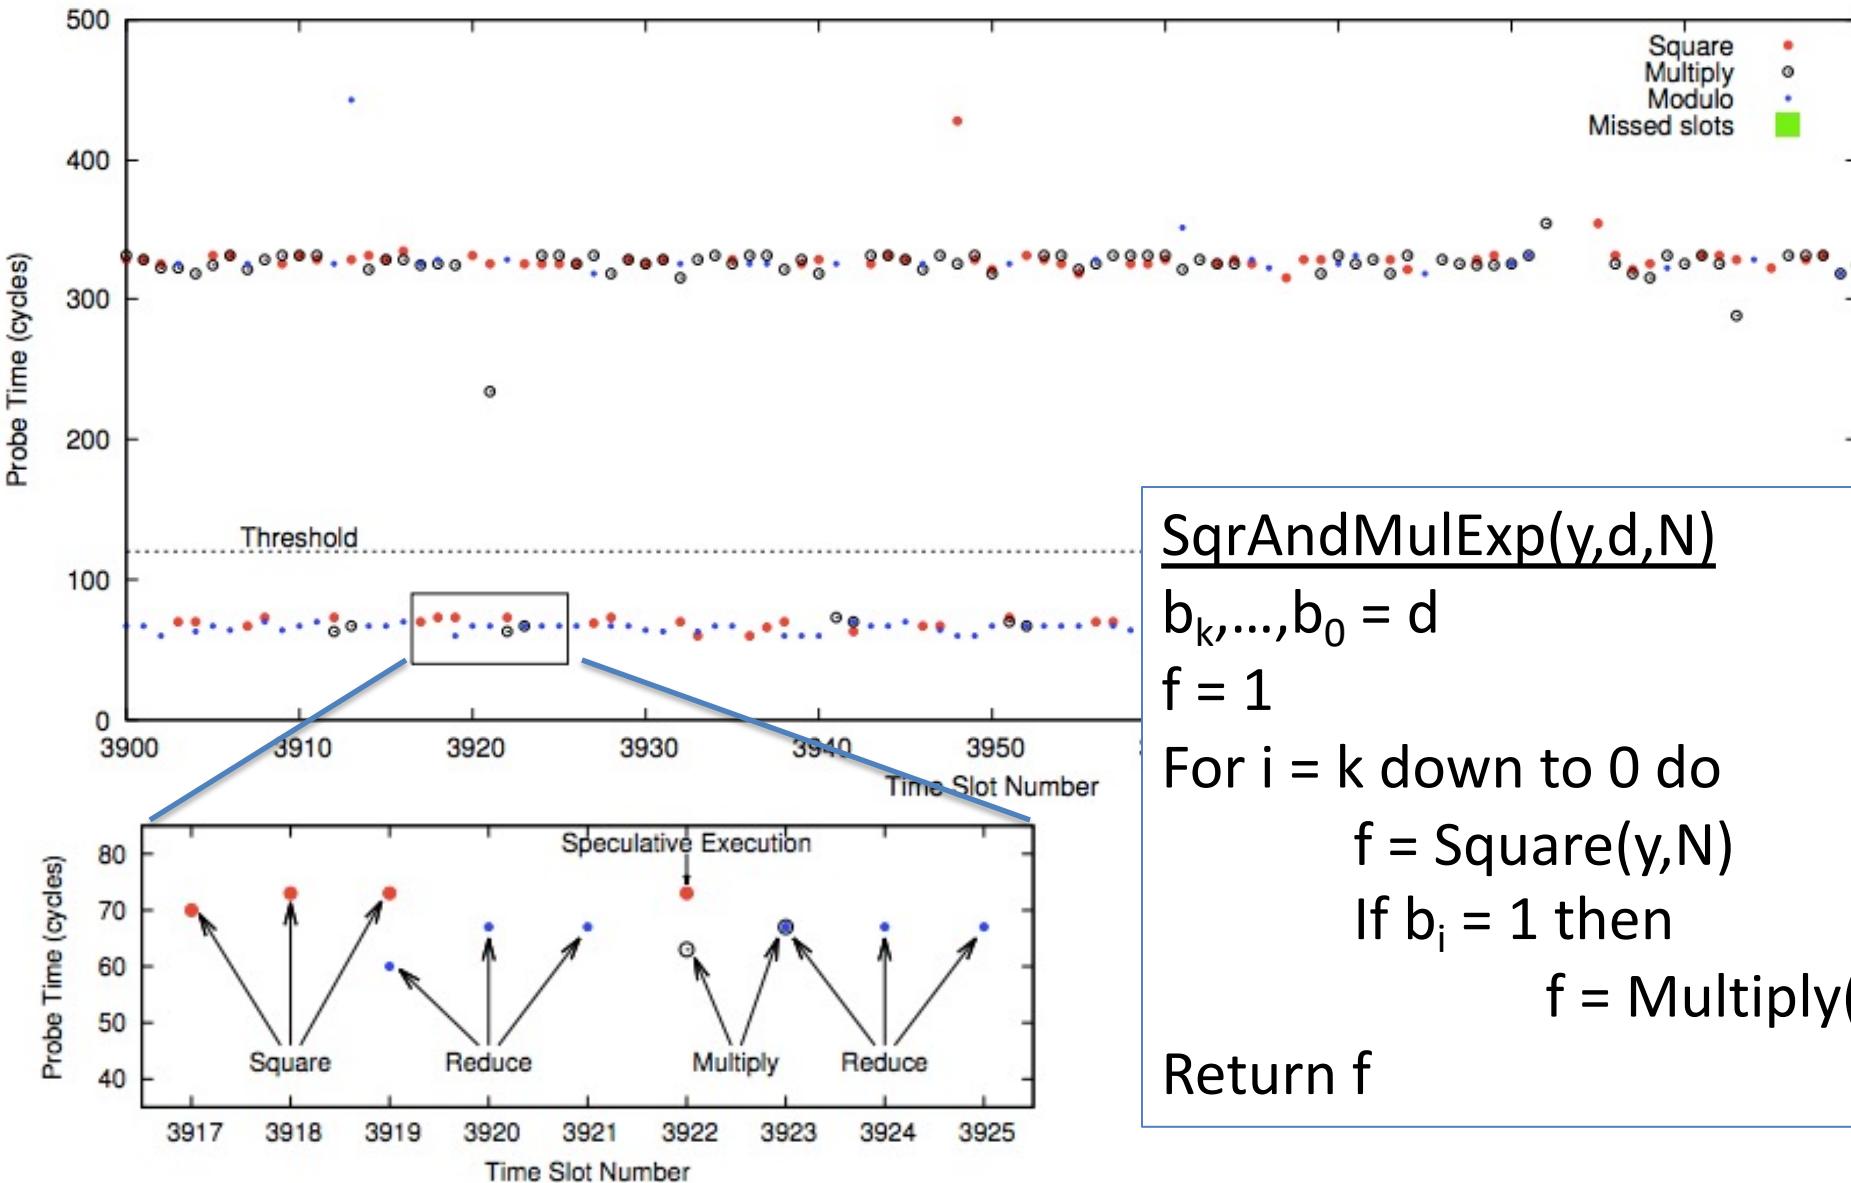
Flush+Reload protocol



Flush+Reload protocol



Attacking Square and Multiply



SqrAndMulExp(y, d, N)

$$b_k, \dots, b_0 = d$$

$$f = 1$$

For $i = k$ down to 0 do

$$f = \text{Square}(y, N)$$

If $b_i = 1$ then

$$f = \text{Multiply}(f, y, N)$$

Return f

Flush+Reload applicability

- Useful anywhere code is shared across processes / VMs / containers. Memory dedup or shared libraries
 - Deduplication turned off in Amazon EC2, but available in modern hypervisors
 - Different VMs do not share libraries
- Cross-tenant attacks in platform-as-a-service (PaaS) clouds
 - [Zhang et al. 2014]
- Used as building block for Meltdown, Spectre

Meltdown



- Discovered independently by Google Project Zero, Cyberus Technology, Graz University
- Speculative execution bugs in Intel x86, ARM, IBM processors + cache-based side-channels (Flush+Reload)
 - Allows reading kernel (or hypervisor, other VM) memory

Intel didn't warn US government about CPU security flaws until they were public

Meltdown and Spectre were kept secret

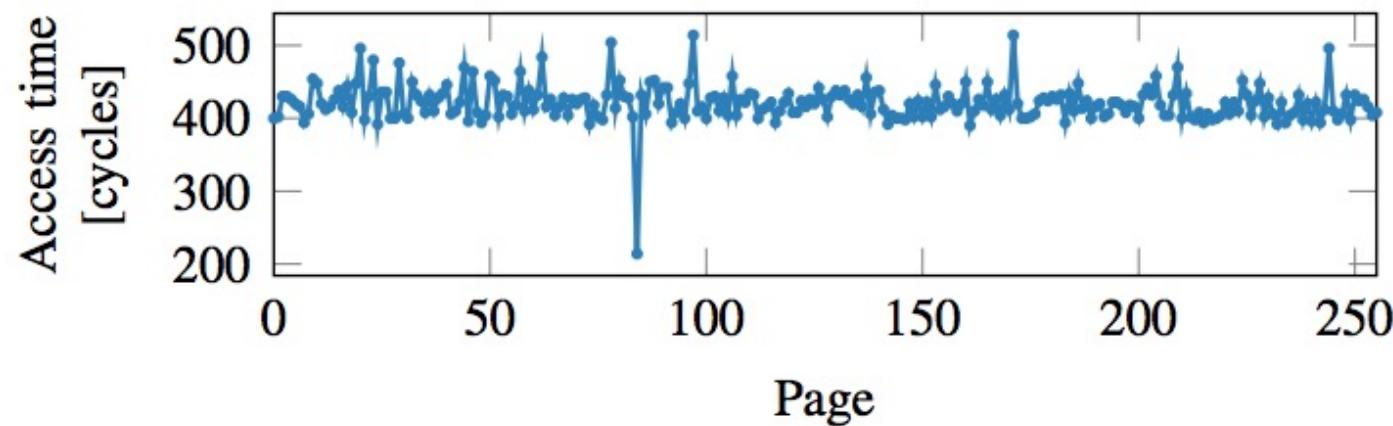
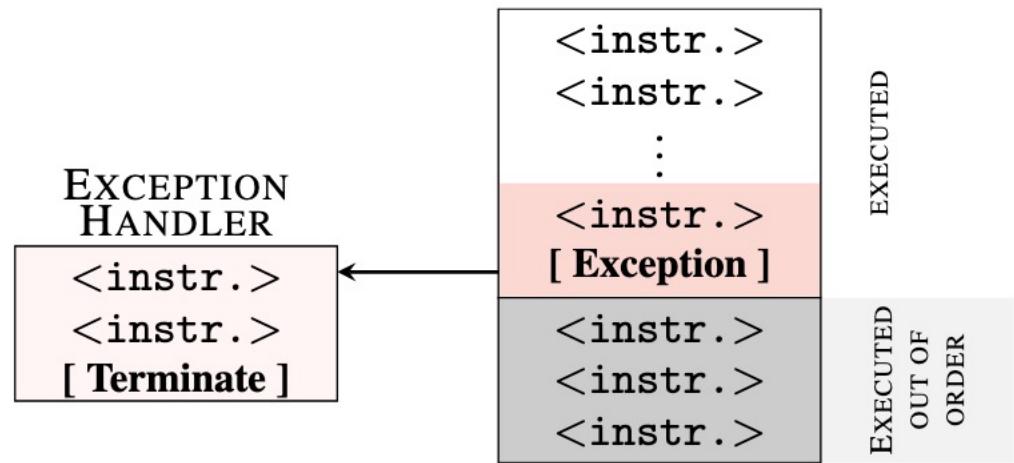
Researchers find malware samples that exploit Meltdown and Spectre

As of Feb. 1, antivirus testing firm AV-TEST had found 139 malware samples that exploit Meltdown and Spectre. Most are not very functional, but that could change.

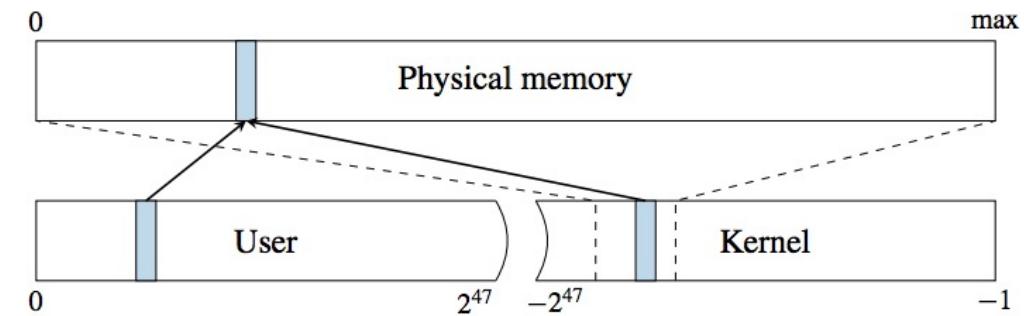
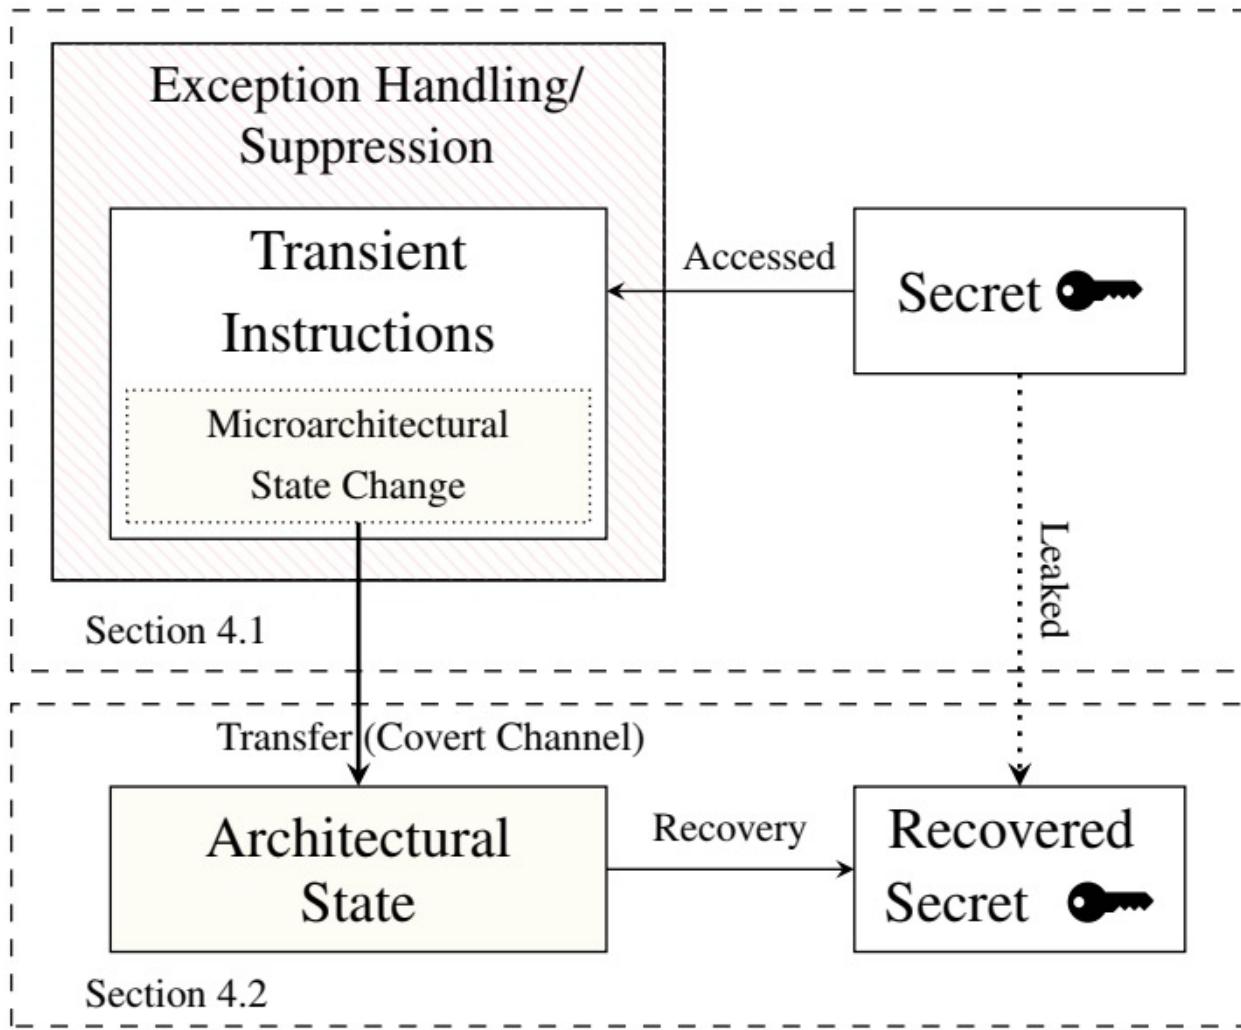
Meltdown: intuition

```
1 raise_exception();
2 // the line below is never reached
3 access(probe_array[data * 4096]);
```

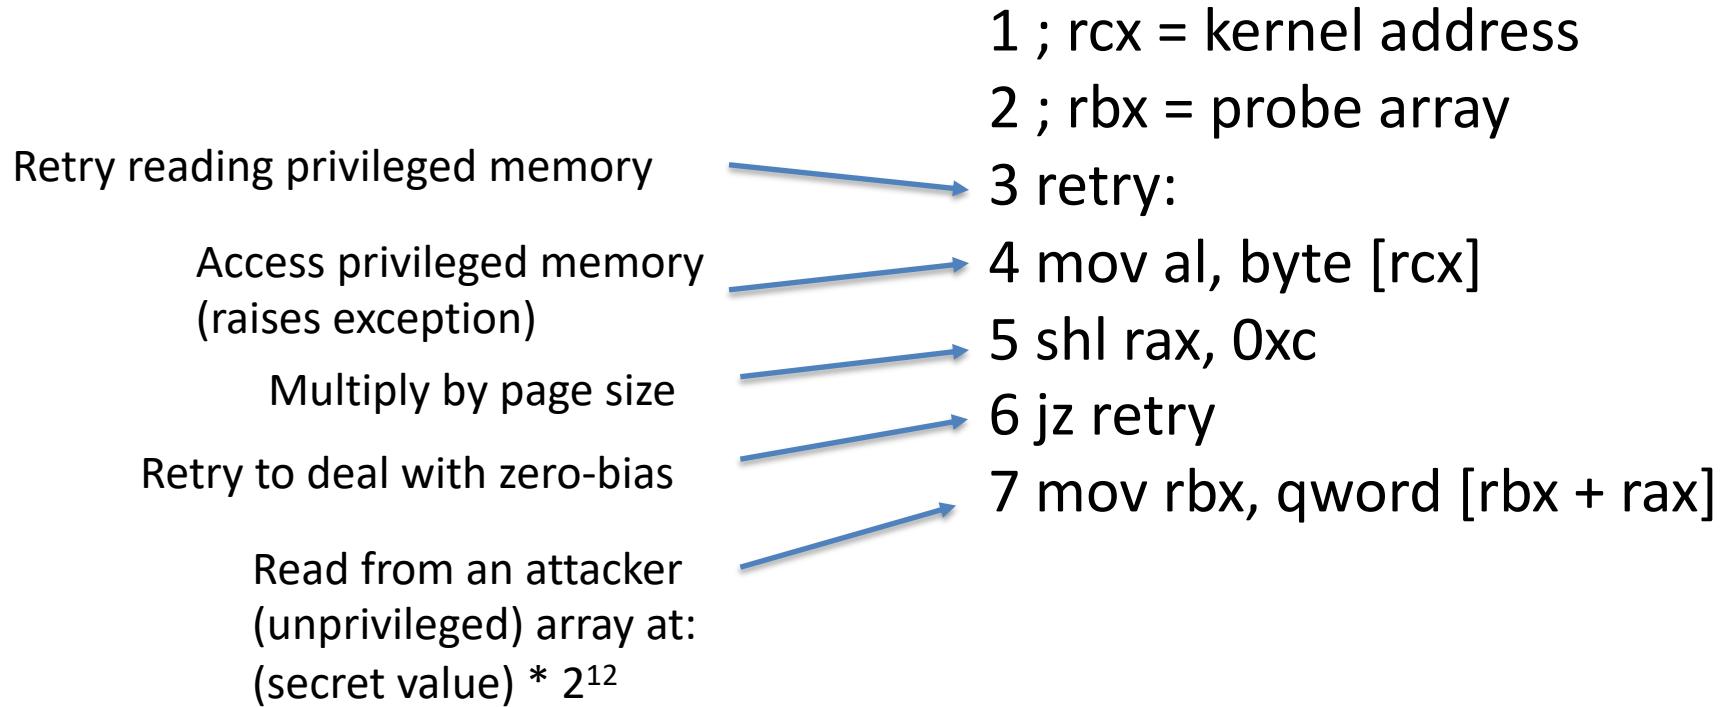
Ensures each value of data triggers
access to distinct memory pages



Meltdown: design



Meltdown: core spy code



Attacker times accessing $[rbx + rax]$ for different values of rax
When finds one that loads fast, learns sensitive byte

Meltdown: nuances

- Privileged memory access raises segfault
 - Would normally crash process
 - Wrap in TXT or include signal handler
- Race condition between segfault and speculative code
 - Keeping speculative code lightweight helps
 - Some errors will creep up, but not many
- Kernel ALSR can be turned on
 - Won't know a priori what to set rcx to
 - Can guess and check (8 GB memory requires ~128 guesses)

Meltdown works very well

```
f94b7690: e5 e5 e5 e5 e5 e5 e5 e5 |.....|  
f94b76a0: e5 e5 e5 e5 e5 e5 e5 e5 |.....|  
f94b76b0: 70 52 b8 6b 96 7f XX XX XX XX XX XX XX XX |pR.k.....|  
f94b76c0: 09 XX |.....|  
f94b76d0: XX |.....|  
f94b76e0: XX 81 |.....|  
f94b76f0: 12 XX e0 81 19 XX e0 81 44 6f 6c 70 68 69 6e 31 |.....Dolphin1|  
f94b7700: 38 e5 e5 e5 e5 e5 e5 e5 |8.....|  
f94b7710: 70 52 b8 6b 96 7f XX XX XX XX XX XX XX |pR.k.....|  
f94b7720: XX |.....|  
f94b7730: XX XX XX XX 4a XX XX XX XX XX XX XX XX |....J.....|  
f94b7740: XX |.....|  
f94b7750: XX e0 81 69 6e 73 74 |.....inst|  
f94b7760: 61 5f 30 32 30 33 e5 e5 e5 e5 e5 e5 e5 |a_0203.....|  
f94b7770: 70 52 18 7d 28 7f XX XX XX XX XX XX XX |pR.}{.....|  
f94b7780: XX |.....|  
f94b7790: XX XX XX XX 54 XX XX XX XX XX XX XX |....T.....|  
f94b77a0: XX |.....|  
f94b77b0: XX 73 65 63 72 |.....secr|  
f94b77c0: 65 74 70 77 64 30 e5 e5 e5 e5 e5 e5 |etpwd0.....|  
f94b77d0: 30 b4 18 7d 28 7f XX XX XX XX XX XX XX |0..}{.....|  
f94b77e0: XX |.....|  
f94b77f0: XX |.....|  
f94b7800: e5 e5 e5 e5 e5 e5 e5 e5 |.....|  
f94b7810: 68 74 74 70 73 3a 2f 2f 61 64 64 6f 6e 73 2e 63 |https://addons.c/  
f94b7820: 64 6e 2e 6d 6f 7a 69 6c 6c 61 2e 6e 65 74 2f 75 |dn.mozilla.net/u/  
f94b7830: 73 65 72 2d 6d 65 64 69 61 2f 61 64 64 6f 6e 5f |ser-media/addon_|  
f94b7840: 69 63 6f 6e 73 2f 33 35 34 2f 33 35 34 33 39 39 |icons/354/354399|  
f94b7850: 2d 36 34 2e 70 6e 67 3f 6d 6f 64 69 66 69 65 64 |-64.png?modified|  
f94b7860: 3d 31 34 35 32 32 34 34 38 31 35 XX XX XX XX XX |=1452244815.....|
```

503 KB/s using TXT “exception suppression”

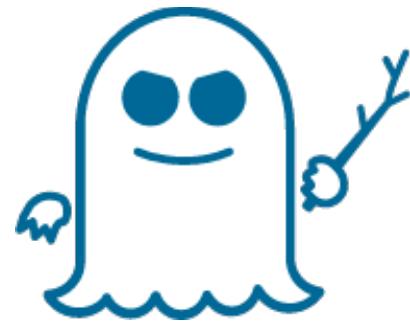
Error rate: 0.02%

Memory dump of Firefox disclosing user passwords

Meltdown: countermeasures

- Minimize kernel memory mapped to virtual address space of process
 - KAISER countermeasure for kernel ASLR already deals with this.
Suggested immediate deployment
- Hard split of virtual memory space into privileged / unprivileged
- Redesign microprocessors
 - For some defenses software microcode updates can be pushed to CPUs

Spectre vulnerability



- Independently discovered by Jann Horn (Google) and Paul Kocher (in collab with some academics)
- Also exploits speculative execution, but differently than Meltdown
 - Doesn't target kernel memory, rather any inaccessible
 - Relies on branch prediction and speculative execution
- KAISER countermeasure doesn't prevent

Microarchitectural vulnerabilities

- Many new attacks after Meltdown/Spectre
 - Foreshadow, Zombieload, Fallout, ...
 - These papers high-impact, widely cited
 - Meltdown paper (2018) has >890 citations
 - Spectre paper (2019) has >1600 citations
- Microarchitectures are complex systems, root of trust in systems. Expect lots of vulnerabilities
 - Vulnerabilities like speculative execution
 - Side-channels, covert channels